

Scaling Fisher Price Architecture with Kubernetes



Jason Carter

Senior Software Engineer

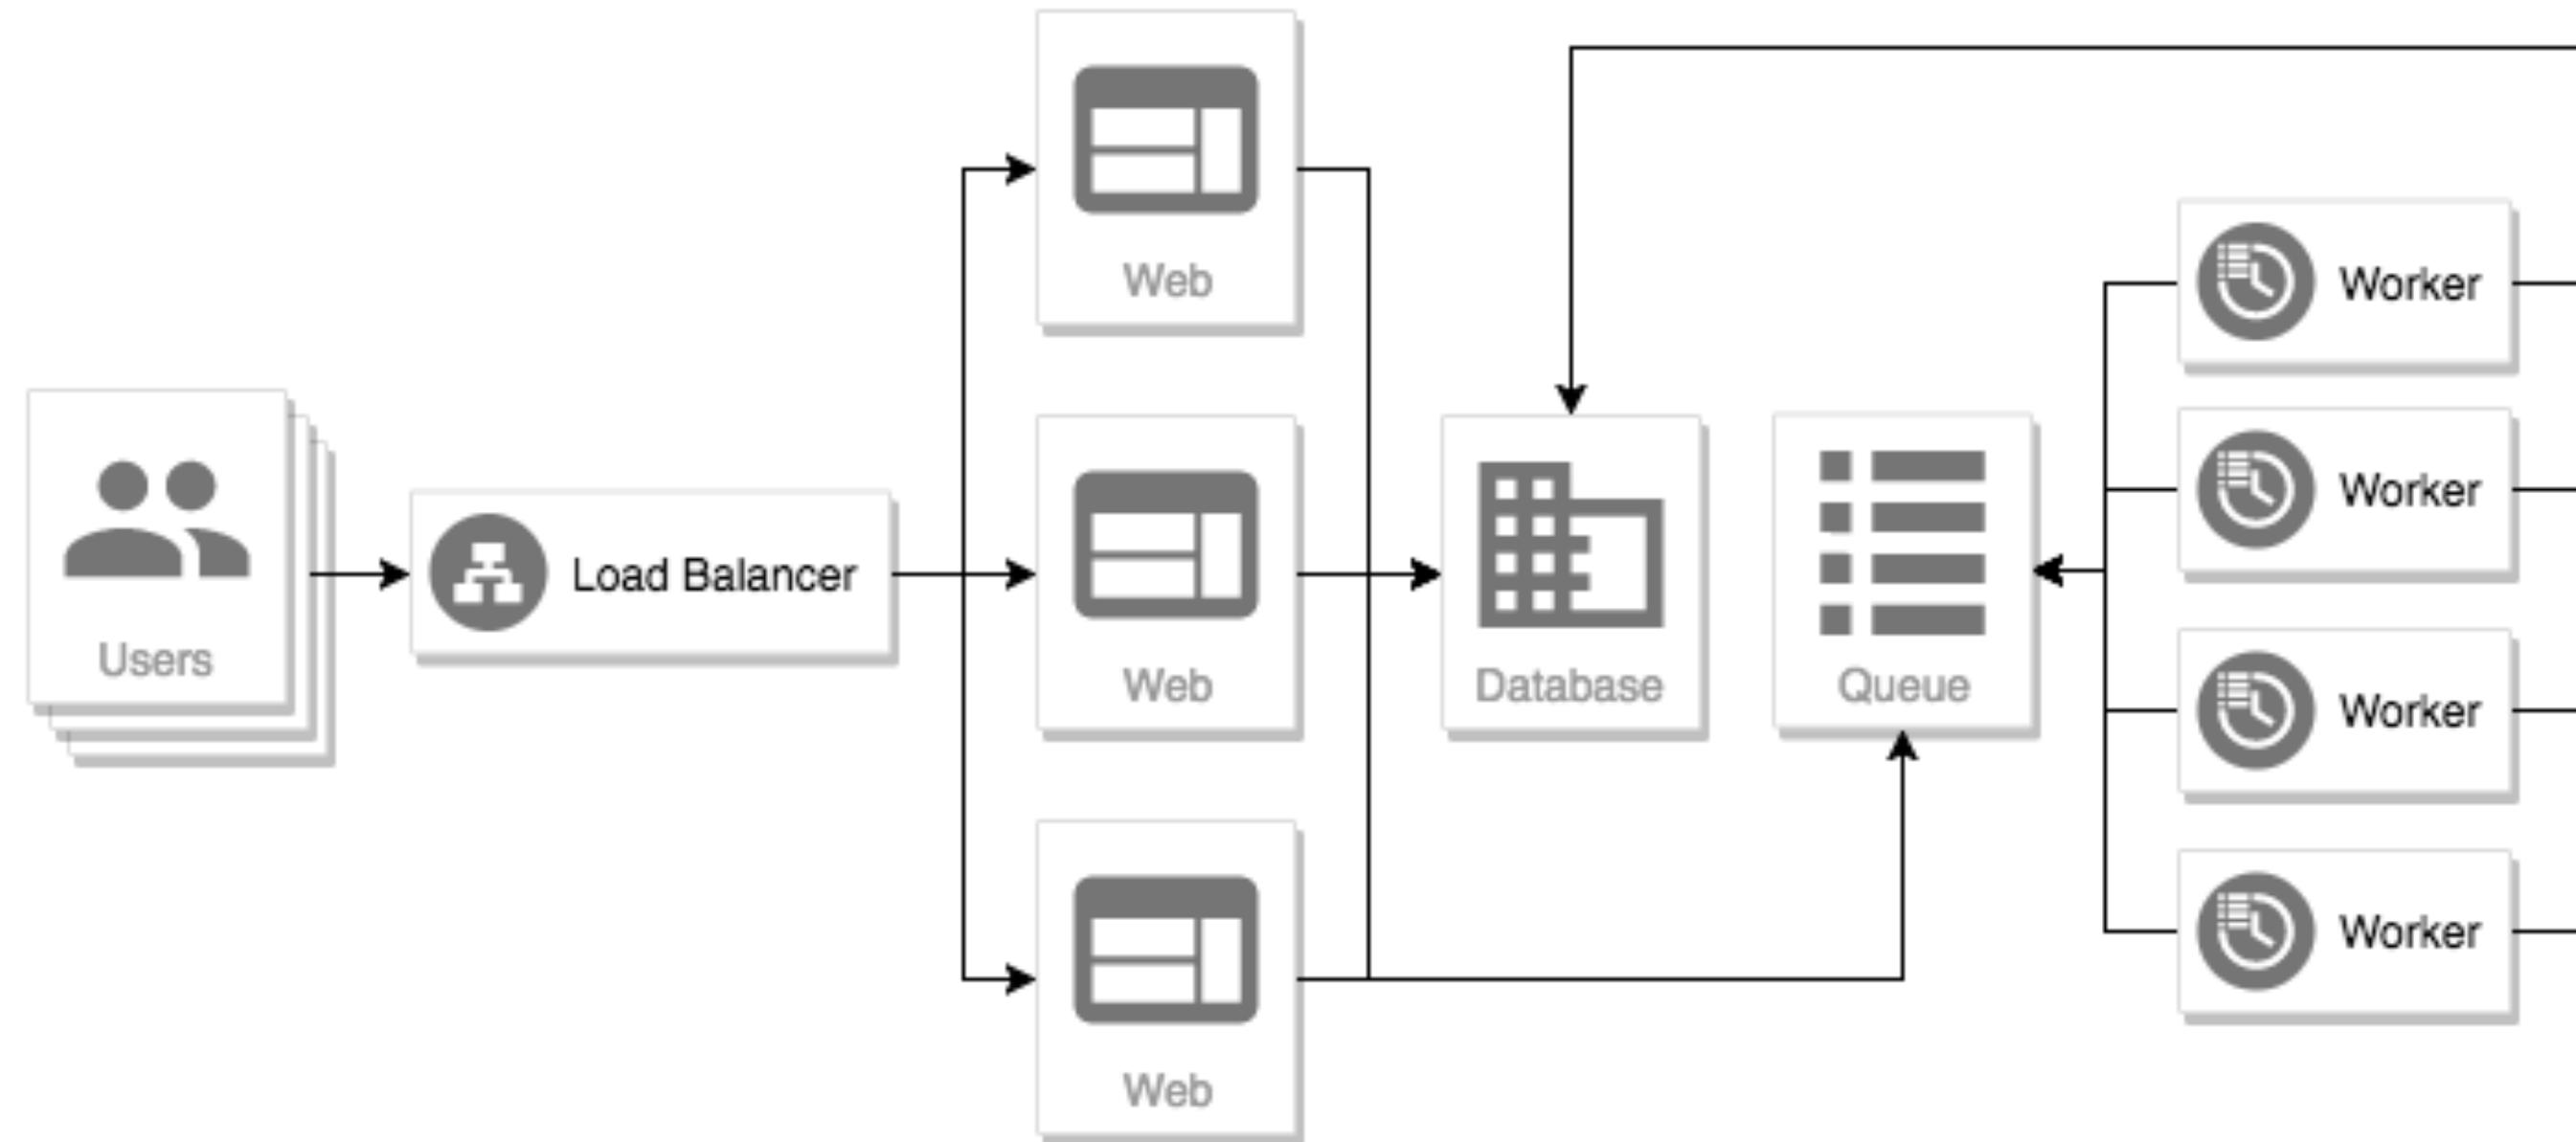


A photograph showing two software engineers from behind, working at their desks. They are both looking at large computer monitors displaying complex code and data visualizations. The engineer on the left is wearing a black t-shirt and has a mug on his desk. The engineer on the right is wearing a green camouflage t-shirt. The monitor on the right has the name "WALTHER" visible in the top right corner.

mavenlink.com/engineering

Mavenlink Integrations

- Connects third party APIs with Mavenlink
- For example,
 - A company might want to use **Jira** for development tasks and **Mavenlink** for accounting, project planning and *client collaboration*.
 - Mavenlink Integrations lets them sync **Jira** and **Mavenlink** so that data flows between both systems.





HEROKU

Pros

- Dead simple to get started
- Easy deploys
- Some monitoring

Cons

- Resource utilization and cost
- Hard to scale queue based workloads
- Less control over the underlying hardware
- Visibility



HEROKU



Why?

- Flexibility
- Better resource usage and scaling
- Cheaper!

Mavenlink ❤ Kubernetes

- Two production applications fully deployed
- Background workers for primary application
- Kubernetes workstations
- Instastage

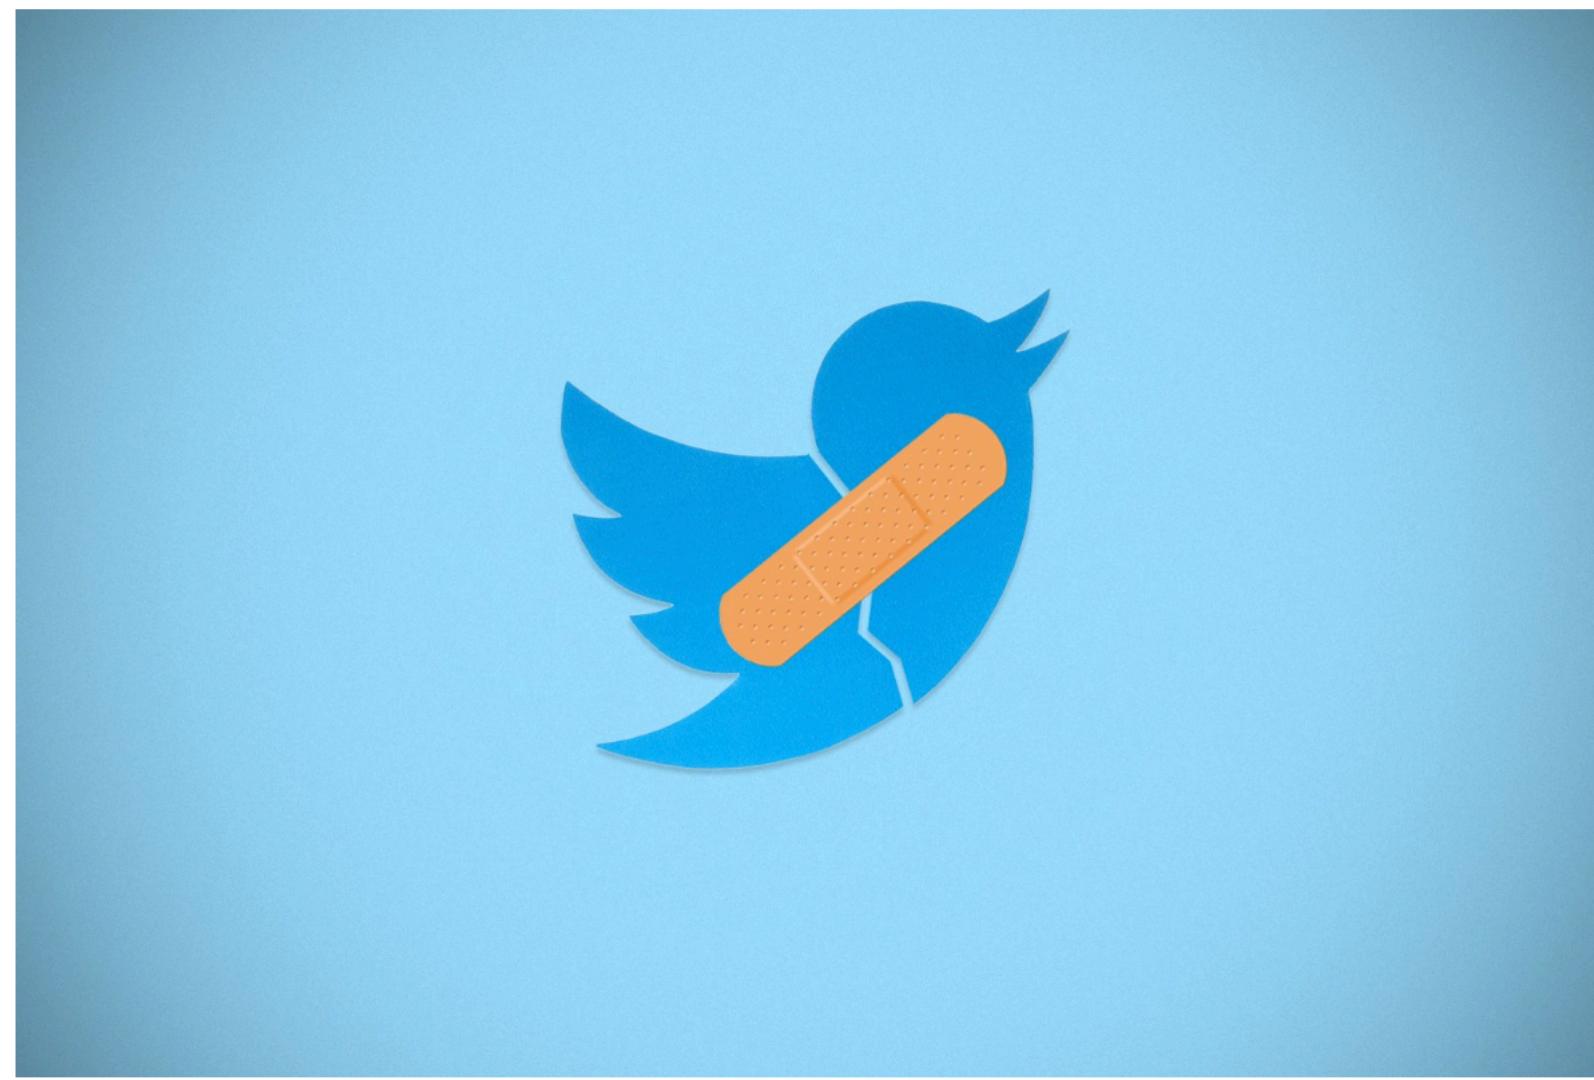
What's the deal with Fisher Price?



“Just an Ass-Backward Tech Company”: How Twitter Lost the Internet War

Twitter faces more challenges than most technology companies: ISIS terrorists, trolls, bots, and Donald Trump. But its last line of defense, the company's head of trust and safety, Del Harvey, isn't making things easier. “Del overcomplicates things . . . and you can see that in the way some of these things are handled publicly.”

BY MAYA KOSOFF | FEBRUARY 19, 2018 6:00 PM





Dare Obasanjo

@Carnage4Life

Follow



Twitter blames their poor handling of abuse on Ruby on Rails? What's next blaming verifying Nazis on MySQL? Way to insult our intelligence.

vanityfair.com/news/2018/02/h...

At the same time, her defenders say, Harvey has been forced to clean up a mess that Twitter should have fixed years ago. Twitter's backend was initially built on Ruby on Rails, a rudimentary web-application framework that made it nearly impossible to find a technical solution to the harassment problem. If Twitter's co-founders had

4:57 PM - 19 Feb 2018

1,010 Retweets 1,912 Likes



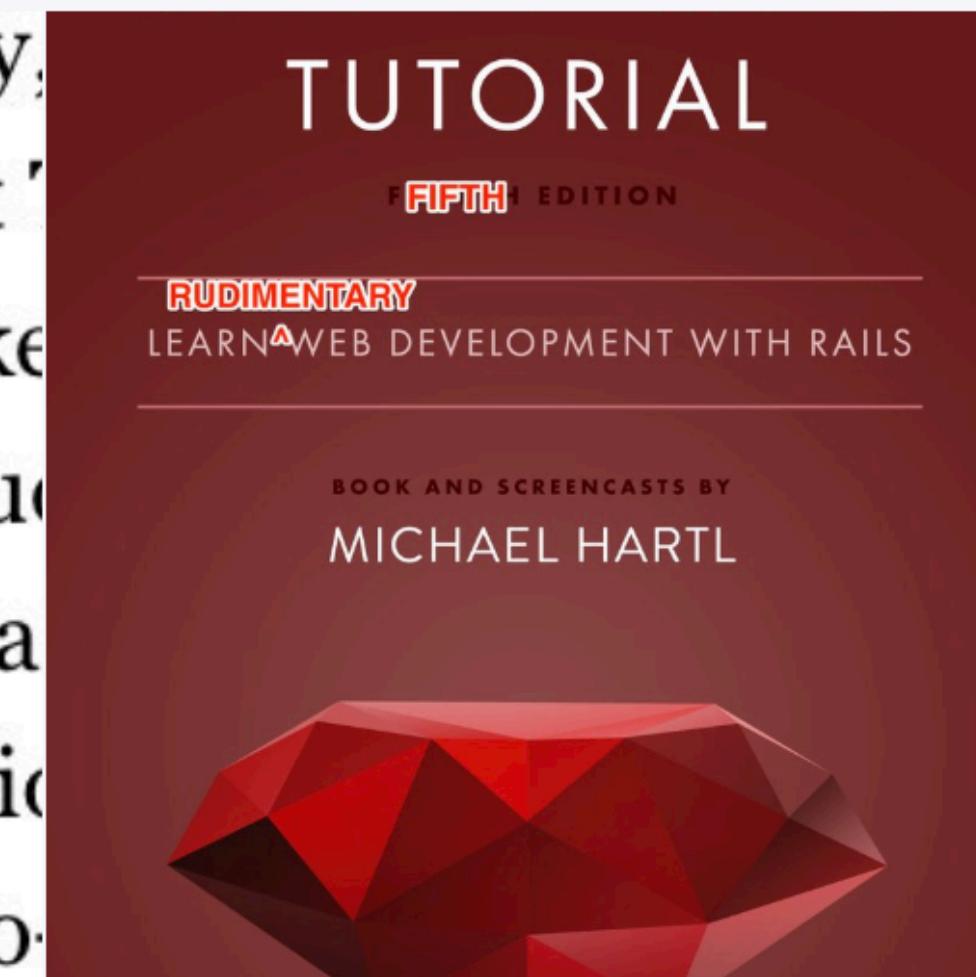


Michael Hartl 🌴 ✅ @mhartl · Feb 19

Replying to @dhh @tobi

OK, the 5th edition of the Ruby on [@RailsTutorial](#) will now be subtitled "Learn Rudimentary Web Development with Rails"

ime, her defenders say, to clean up a mess that's been years ago. Twitter's back on Ruby on Rails, a robust framework that makes it easy to find a technical solution to any problem. If Twitter's co-



2



2



47





DHH

@dhh

Follow



Maybe we can forward the question to [@tobi](#) ? How's the Fisher-Price infrastructure holding back your \$13B public company?

April Joyner @aprjoy

Twitter's failure on this issue is not a joke, but at "blaming verifying Nazis on MySQL."

Also, I wonder what @dhh thinks of the claim that @rails is "Fisher-Price..."

6:58 PM - 19 Feb 2018

75 Retweets 372 Likes





**Rails isn't the
problem...**

**...how we're using
and deploying it
might be.**



What we'll cover

- A traditional "production" Rails application
- What is Kubernetes and how can it help me scale my application?
- Advanced Kubernetes Concepts

A disclaimer



**Lets look at a
traditional Rails
application...**

We'll discuss

- Developing
- Shipping it to *production*

**Let's identify some
common scaling
problems**

Developing

Setting up your development environment

Probably went something like this...

- Install packages
- Install languages, frameworks, and tools
- Install dependencies
- Install services (mySQL, Postgres, Redis, ElasticSearch, etc...)

**what about
replicating this to all
of your other
developers
machines?**

Problem 1

Managing an applications dependencies across development machines is difficult

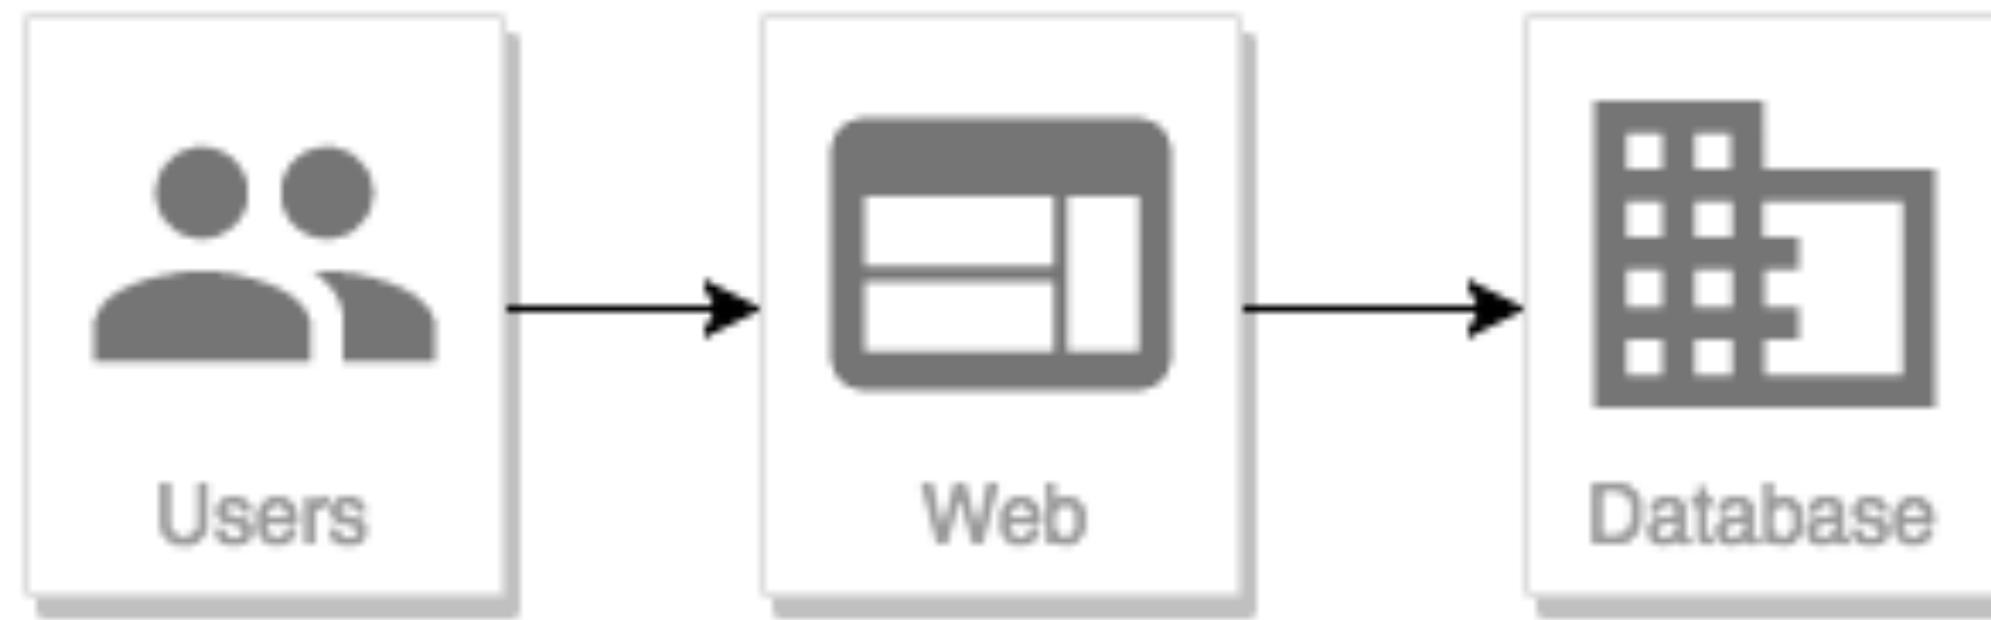
Shipping it to production

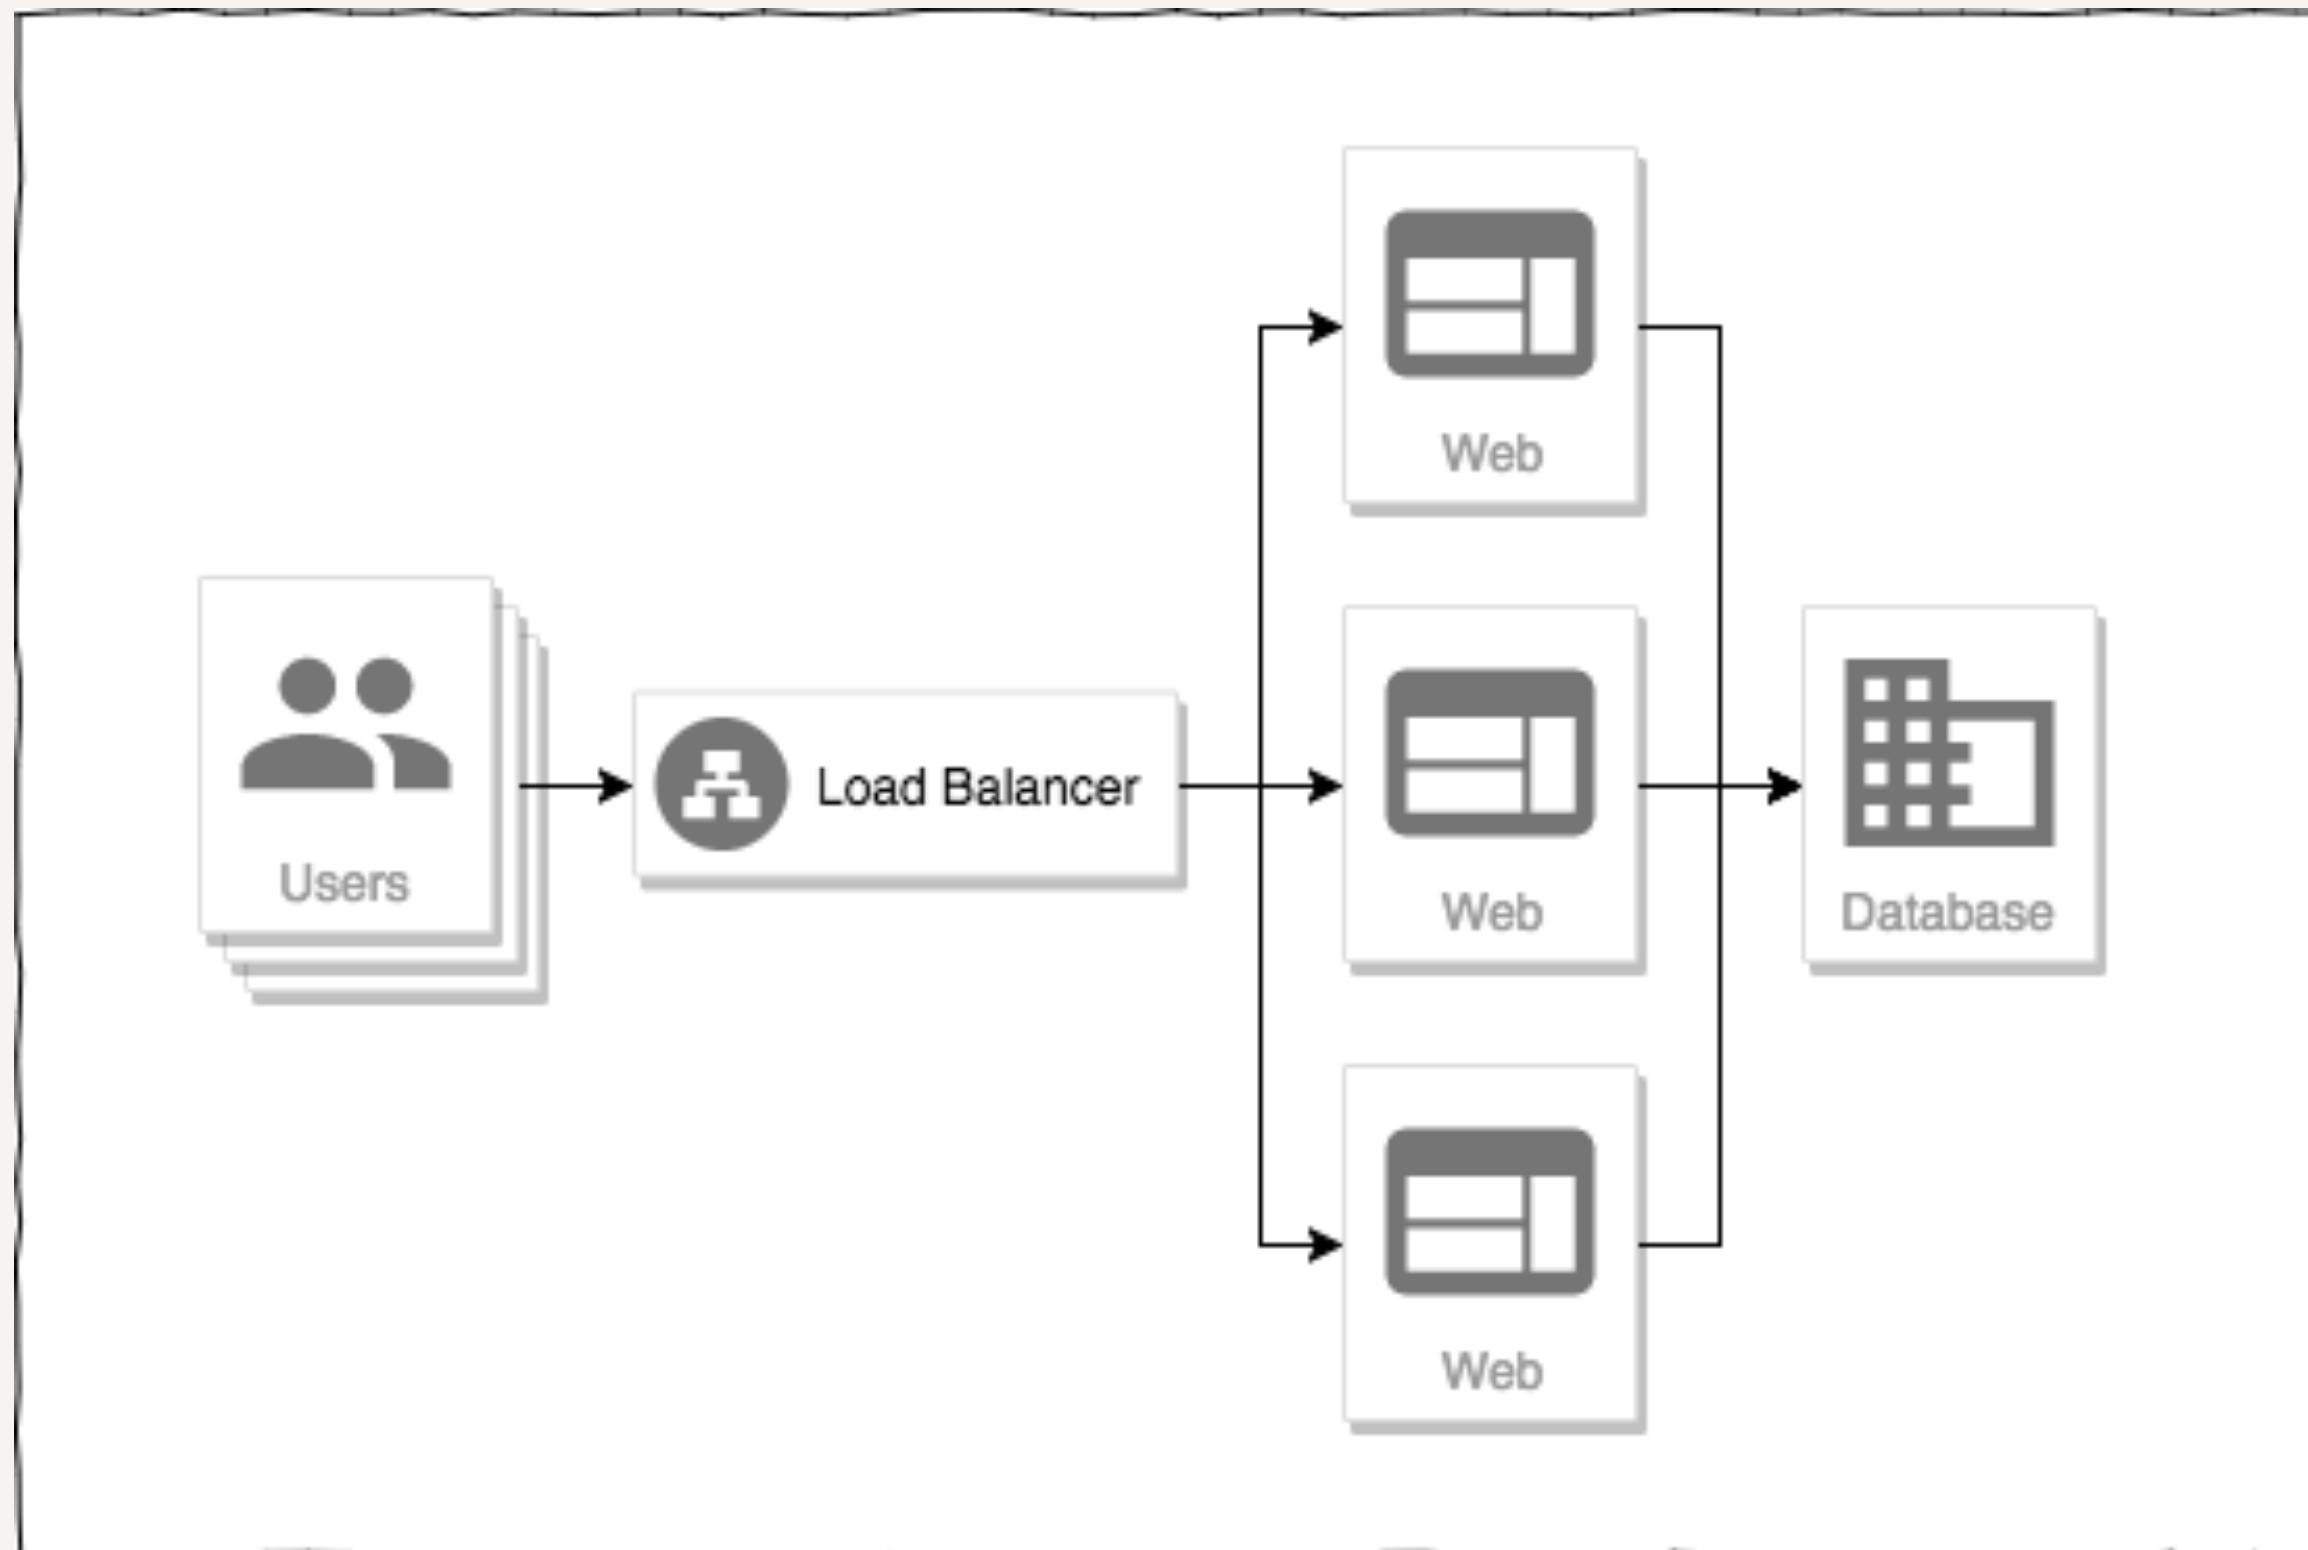
- Need a database
- Need some servers to run my application

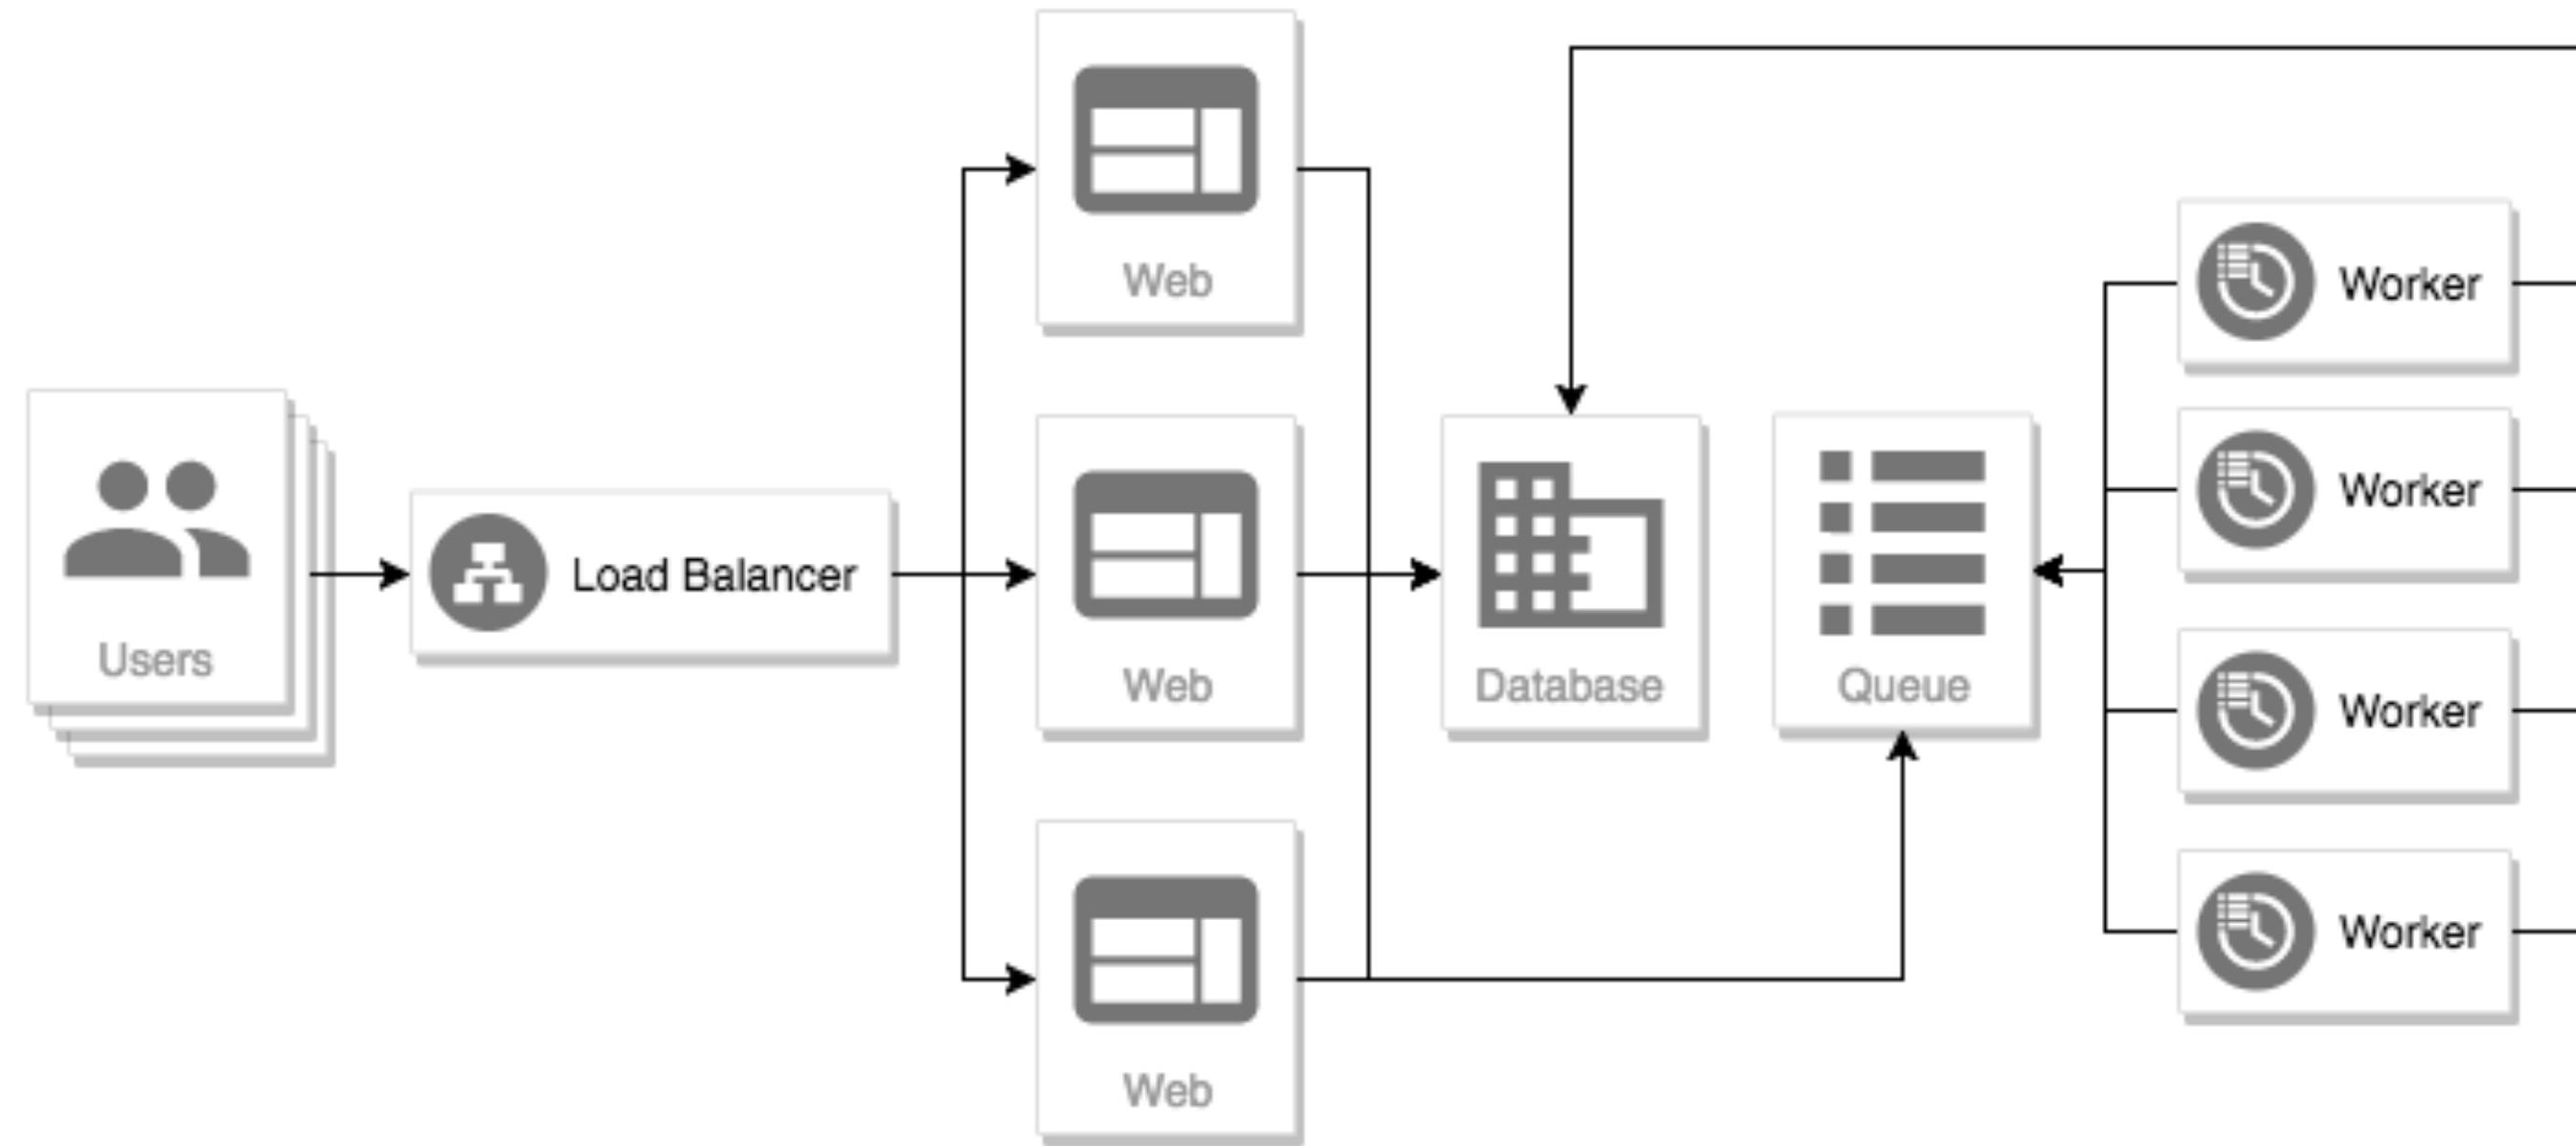
Options

- Roll your own servers
- Use Heroku
- Use a cloud service provider

**Lets go with a cloud
provider for this
example...**







This quickly becomes quite the chore!

- Maintaining server inventory
- Maintaining deploy and development tooling
- Monitoring
- Rollbacks
- Backups

More problems!

Problem 2

Keeping development, staging, and production environments the same

Problem 3

Scaling inventory with demand

Problem 4

Utilization and cost

Problem 5

What if I need to deploy and manage more services?

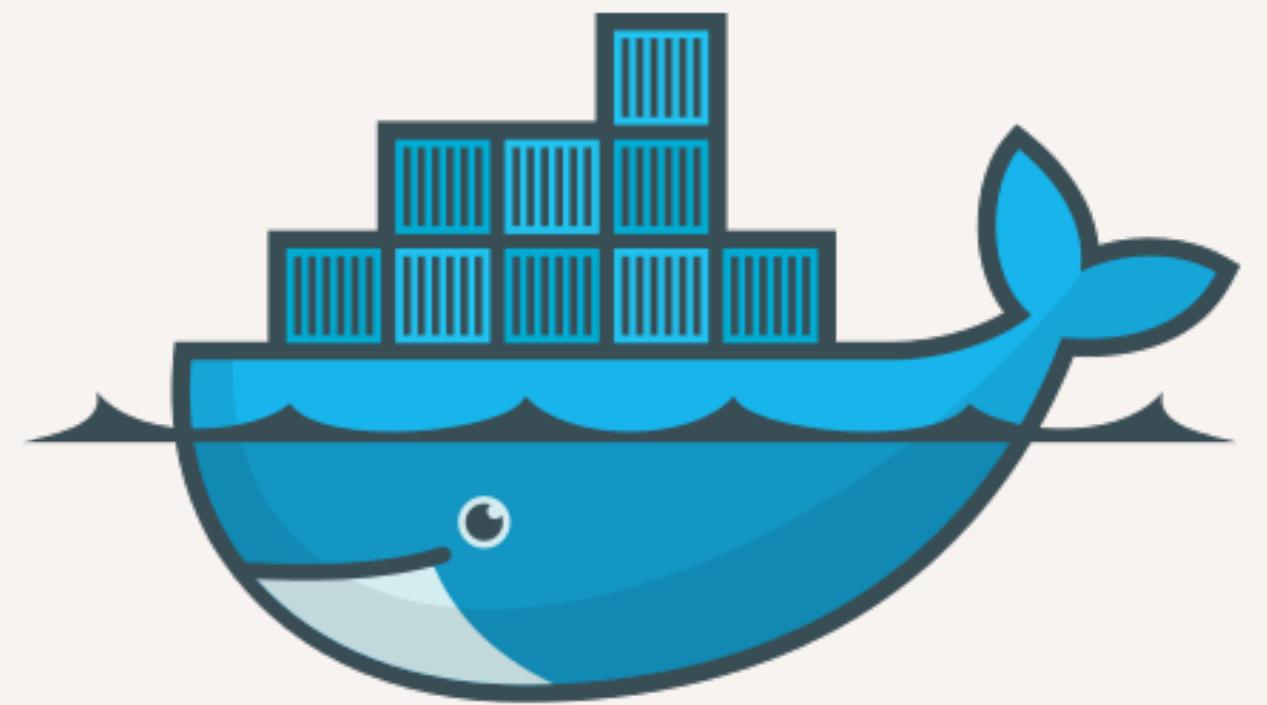
**There has to be a
better way**

**Lets start by
discussing containers**

What's a container anyway?

A container image is a *lightweight*, **stand-alone**, executable package of a piece of software that includes everything needed to run it.

**Containers *run anywhere* and
have everything they need to run**



docker

This solves #1 and #2

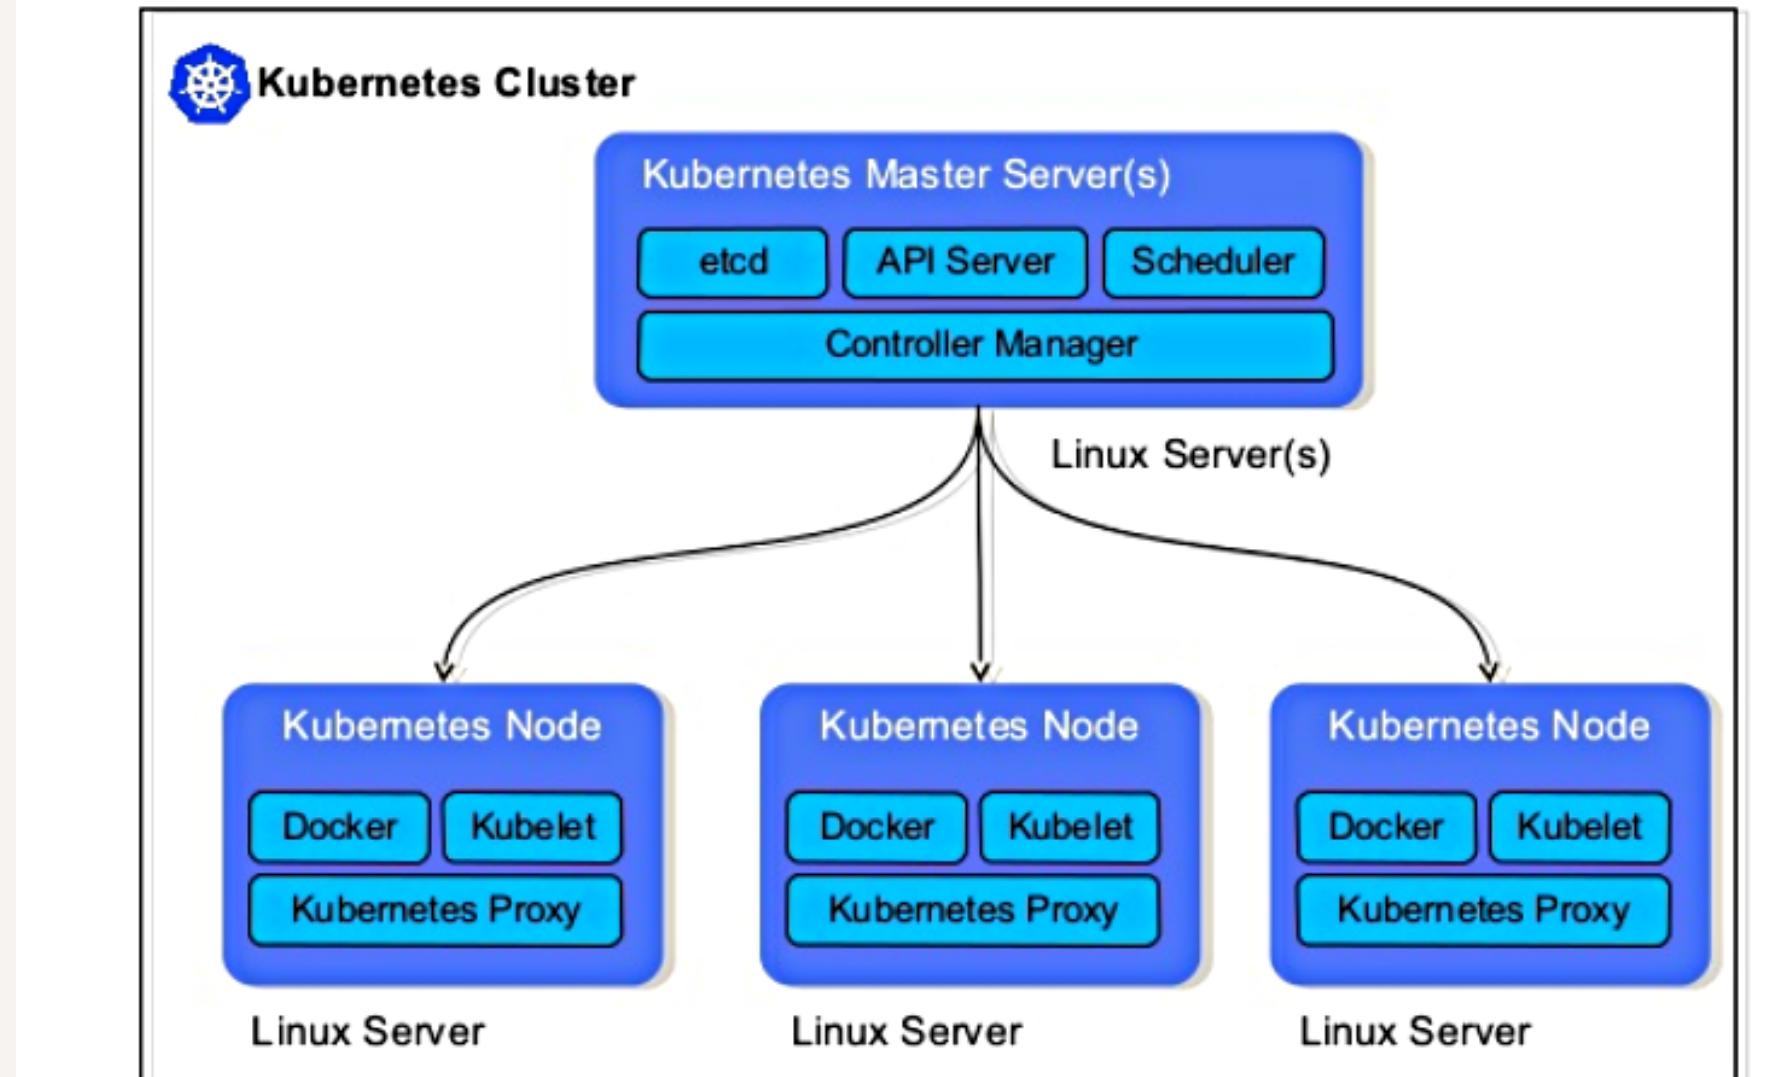
**But now that I've got
a container, how do I
deploy and manage
them?**

What is Kubernetes?

Kubernetes is a portable, extensible open-source platform for managing containerized workloads and services, that facilitates both declarative configuration and automation¹.

¹ [What is Kubernetes?, kubernetes.io.](https://kubernetes.io)

Kubernetes Architectural Overview

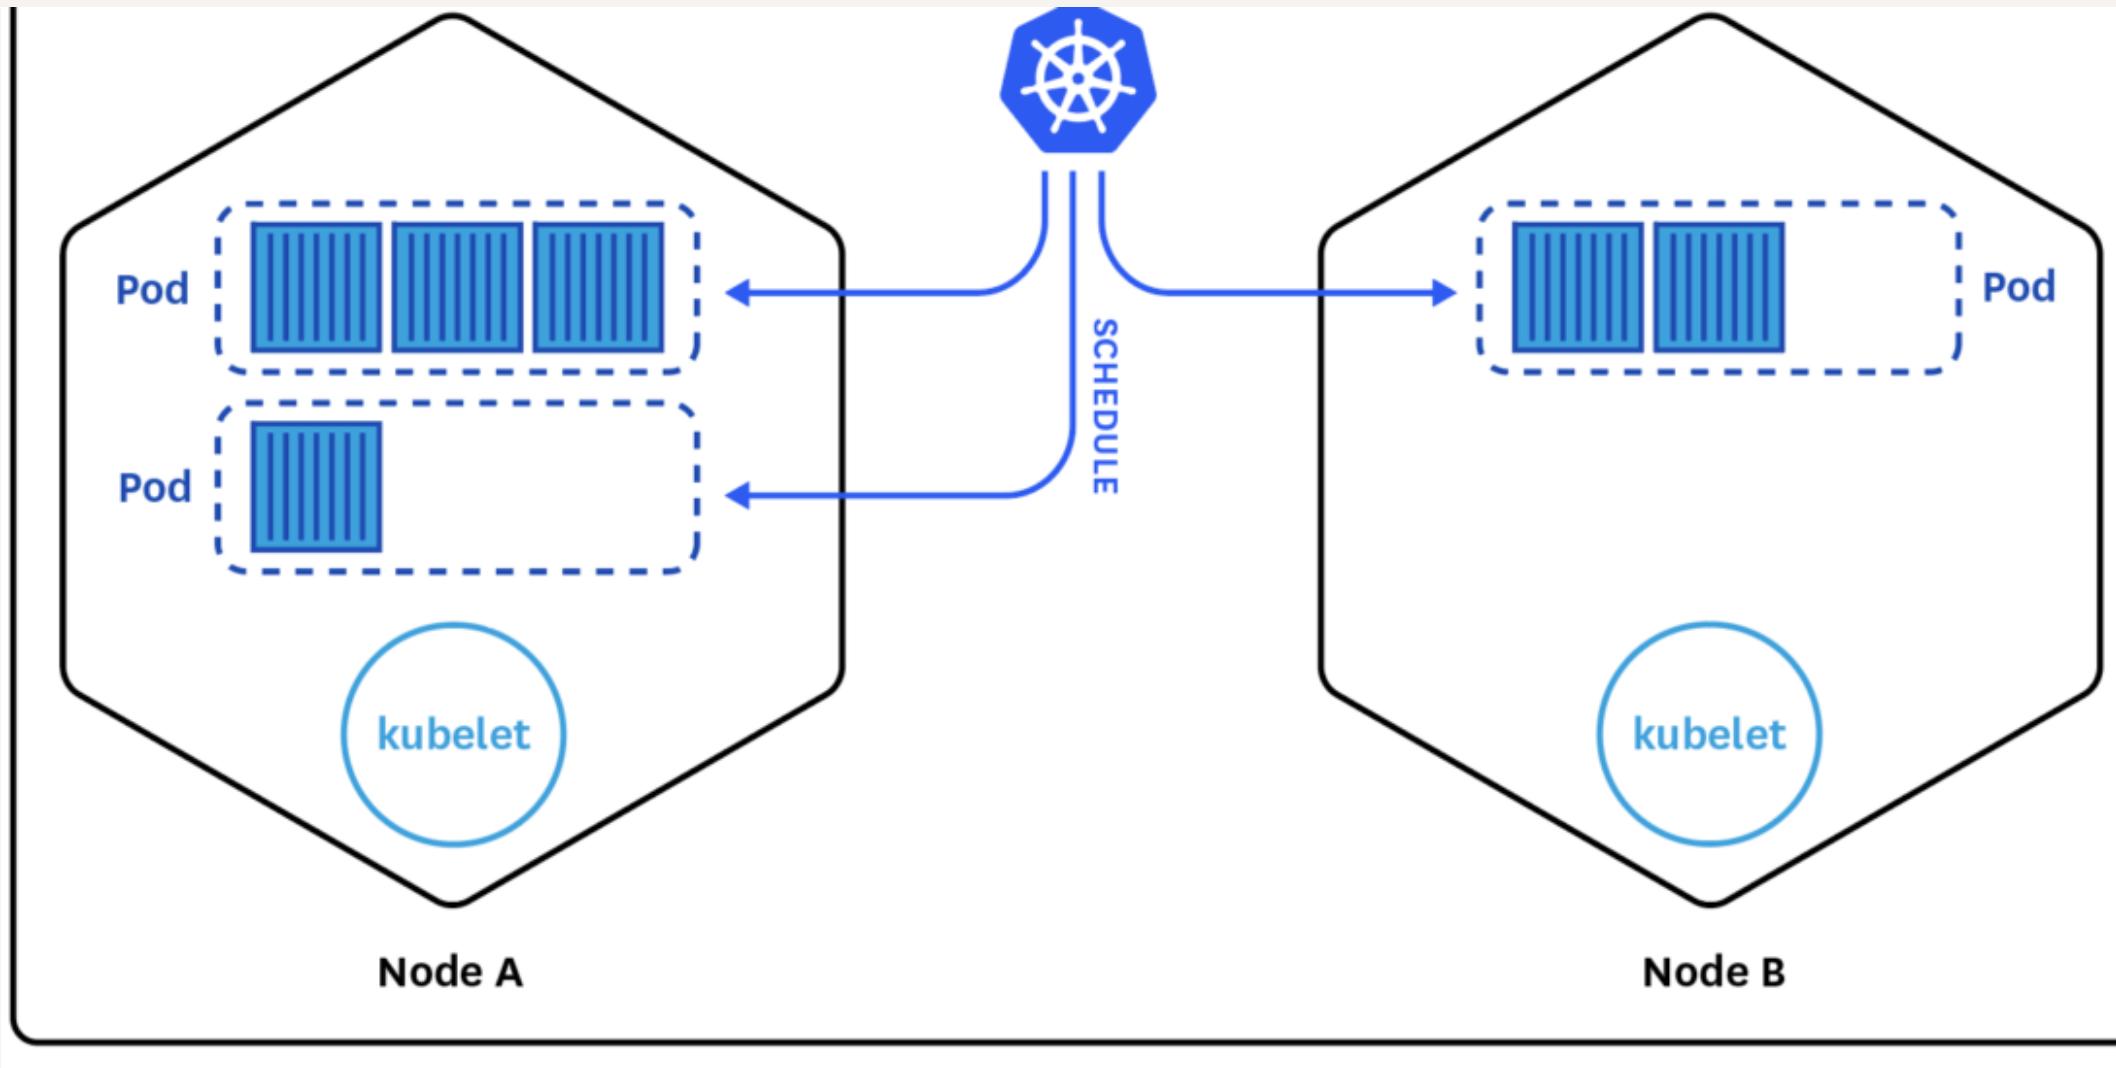


@wattsteve



Kubernetes Architectural Overview (by Steve Watt, Red Hat).

Scheduling



Unsourced 😊

What do I like about Kubernetes?

- Abstractions for thinking about infrastructure
- Don't have to know **too** much about the underlying hardware
- Trivial scaling
- Very, very, flexible!

**How do I get my Rails
application on
Kubernetes?**

1) Spin up a cluster!



Microsoft Azure



Google Cloud Platform

2) Containerize it!

```
FROM ruby:2.5

USER root
RUN apt-get update && apt-get install -qq -y build-essential nodejs libpq-dev --fix-missing --no-install-recommends

RUN mkdir -p /app
WORKDIR /app

COPY Gemfile Gemfile
COPY Gemfile.lock Gemfile.lock

USER app
RUN RAILS_ENV=production bundle install --jobs=4 --retry=3 --deployment

ENV RAILS_ENV=production

COPY . .

USER root
RUN chown -R app. /app

USER app
CMD ["bundle", "exec", "rails s"]
```

3) Put it on Kubernetes

**We're going to write some
manifests that describe what our
application should look like**

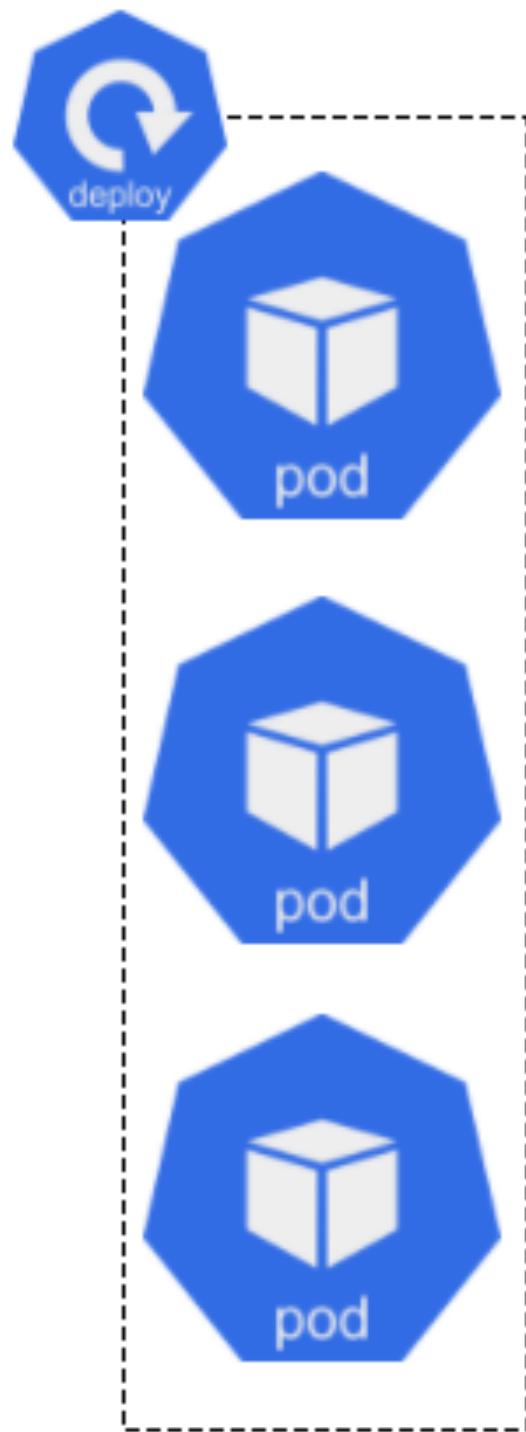
Deployment

A deployment declares your desired state of an application, usually as a series of pods and how they should be managed

Pod

A pod (as in a pod of whales or pea pod) is a group of one or more containers (such as Docker containers), with shared storage/network, and a specification for how to run the containers.

In laymans terms, pods are units of works or replicas of an application. And deployment describes how to configure them.



```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: rails
spec:
  replicas: 3
  template:
    metadata:
      name: rails-pod
      labels:
        service: rails-web
    spec:
      containers:
        - name: rails-container
          image: rails:latest
  restartPolicy: Always
```

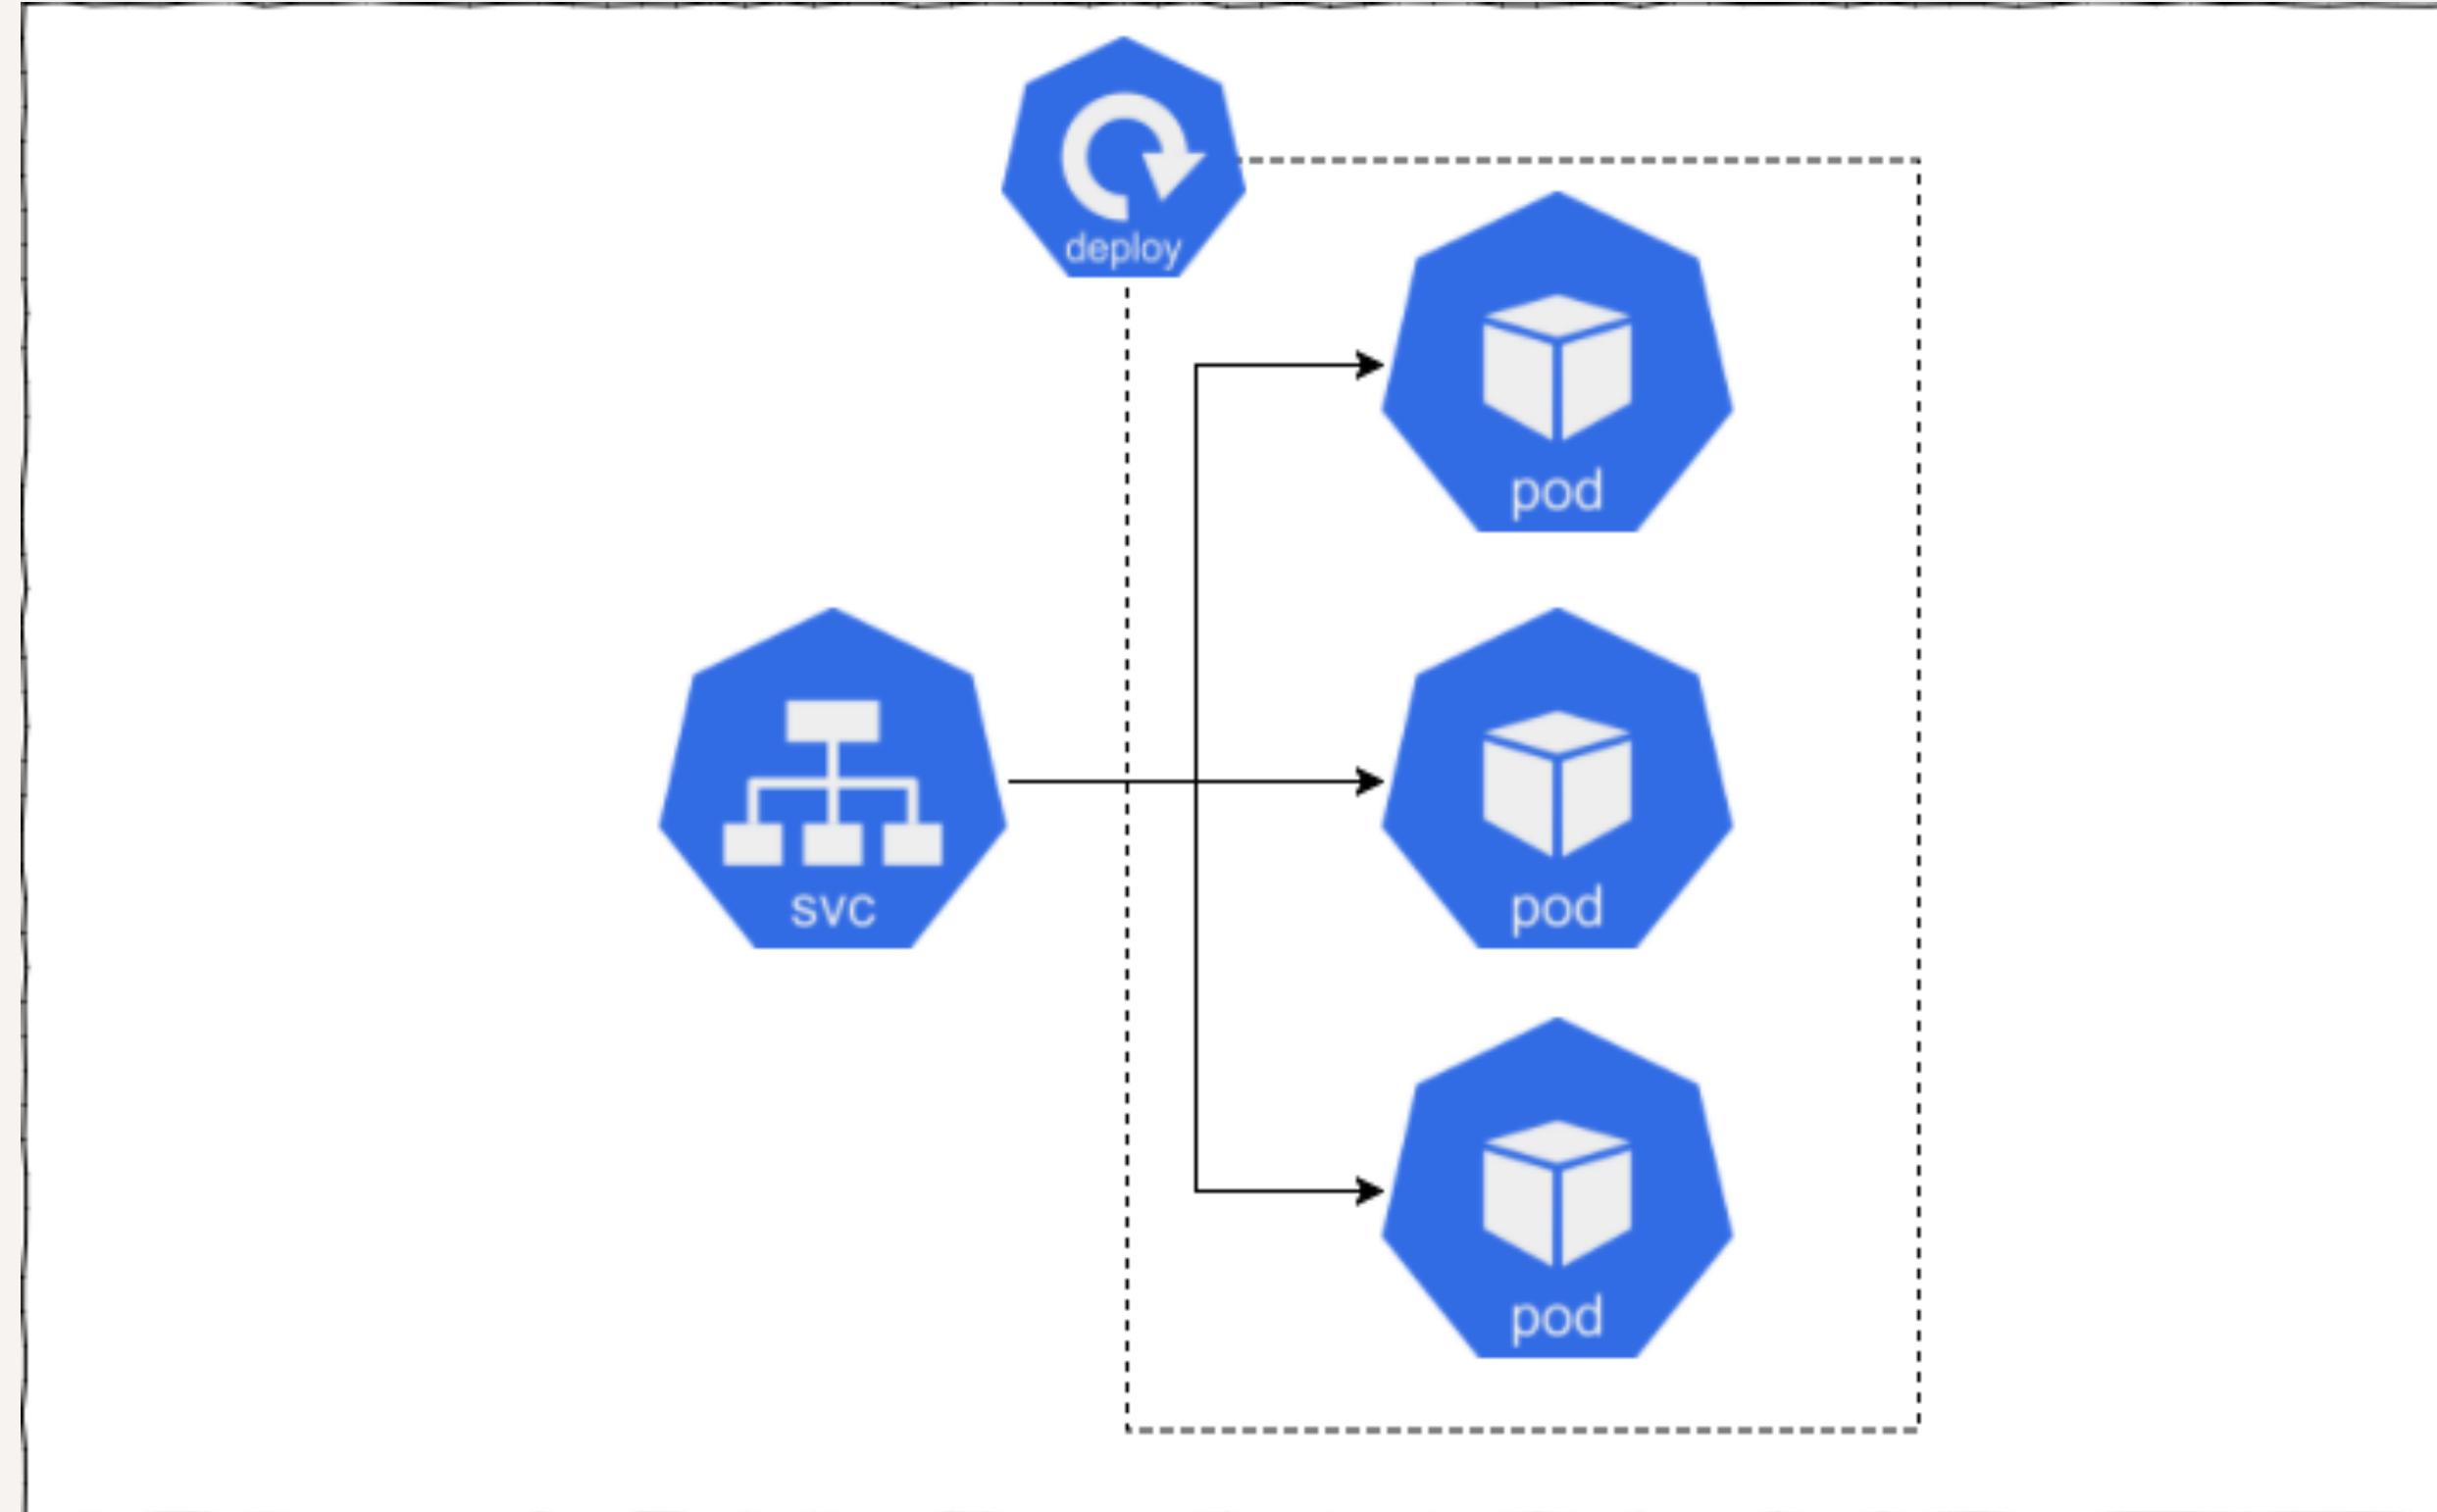
**what can I do with a
deployment?**

```
# How much resources should we allocate to this pod?  
resources:  
    requests:  
        memory: 1Gi  
        cpu: 200m  
    limits:  
        cpu: 1  
        memory: 2Gi
```

```
# Environment variables that the pod should have access to
env:
  - name: MY_FAVORITE_COLOR
    value: blue
  - name: SUPER_TOP_SECRET_USERNAME
    valueFrom:
      secretKeyRef:
        name: secret-token
        key: username
```

Service

A service allows other things to access your deployment with a predictable host.

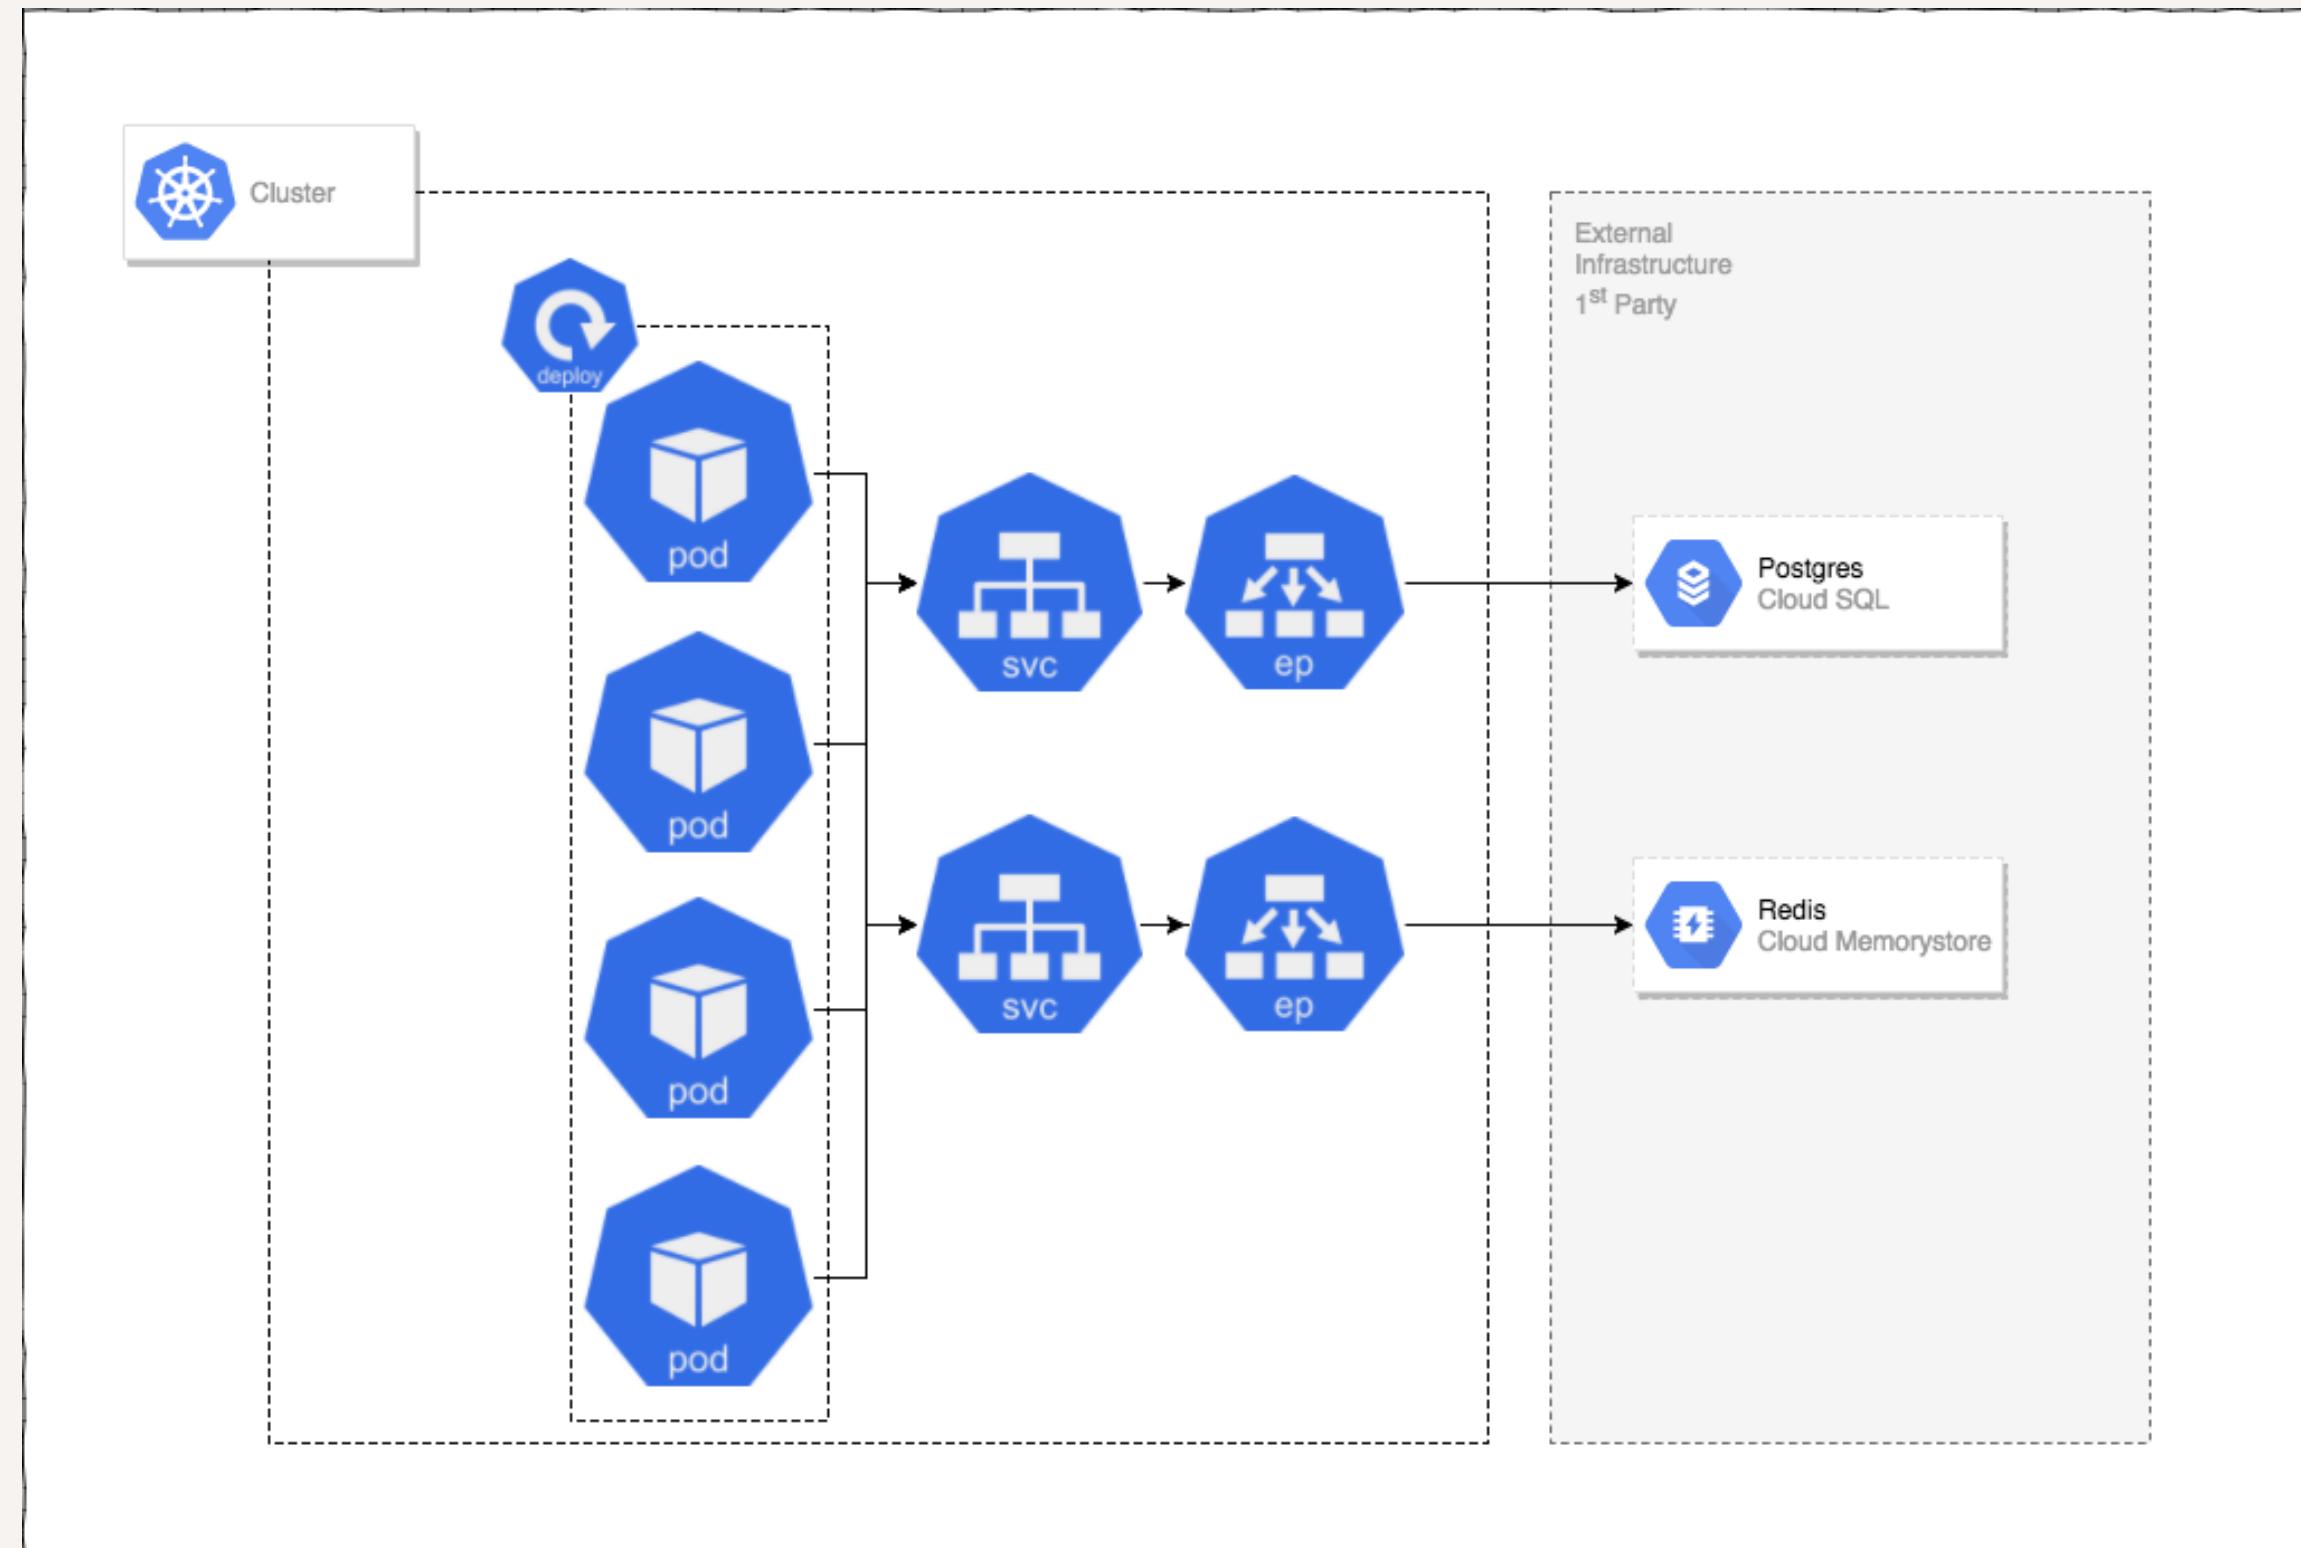


```
apiVersion: v1
kind: Service
metadata:
  labels:
    service: rails
  name: rails-service
spec:
  type: NodePort
  ports:
  - name: rails-port
    port: 3000
    protocol: TCP
    targetPort: 3000
    nodePort: 30000
  selector:
    service: rails-web
```

**This becomes a
human readable (and
predictable) host**

rails-service.default.svc.cluster.local

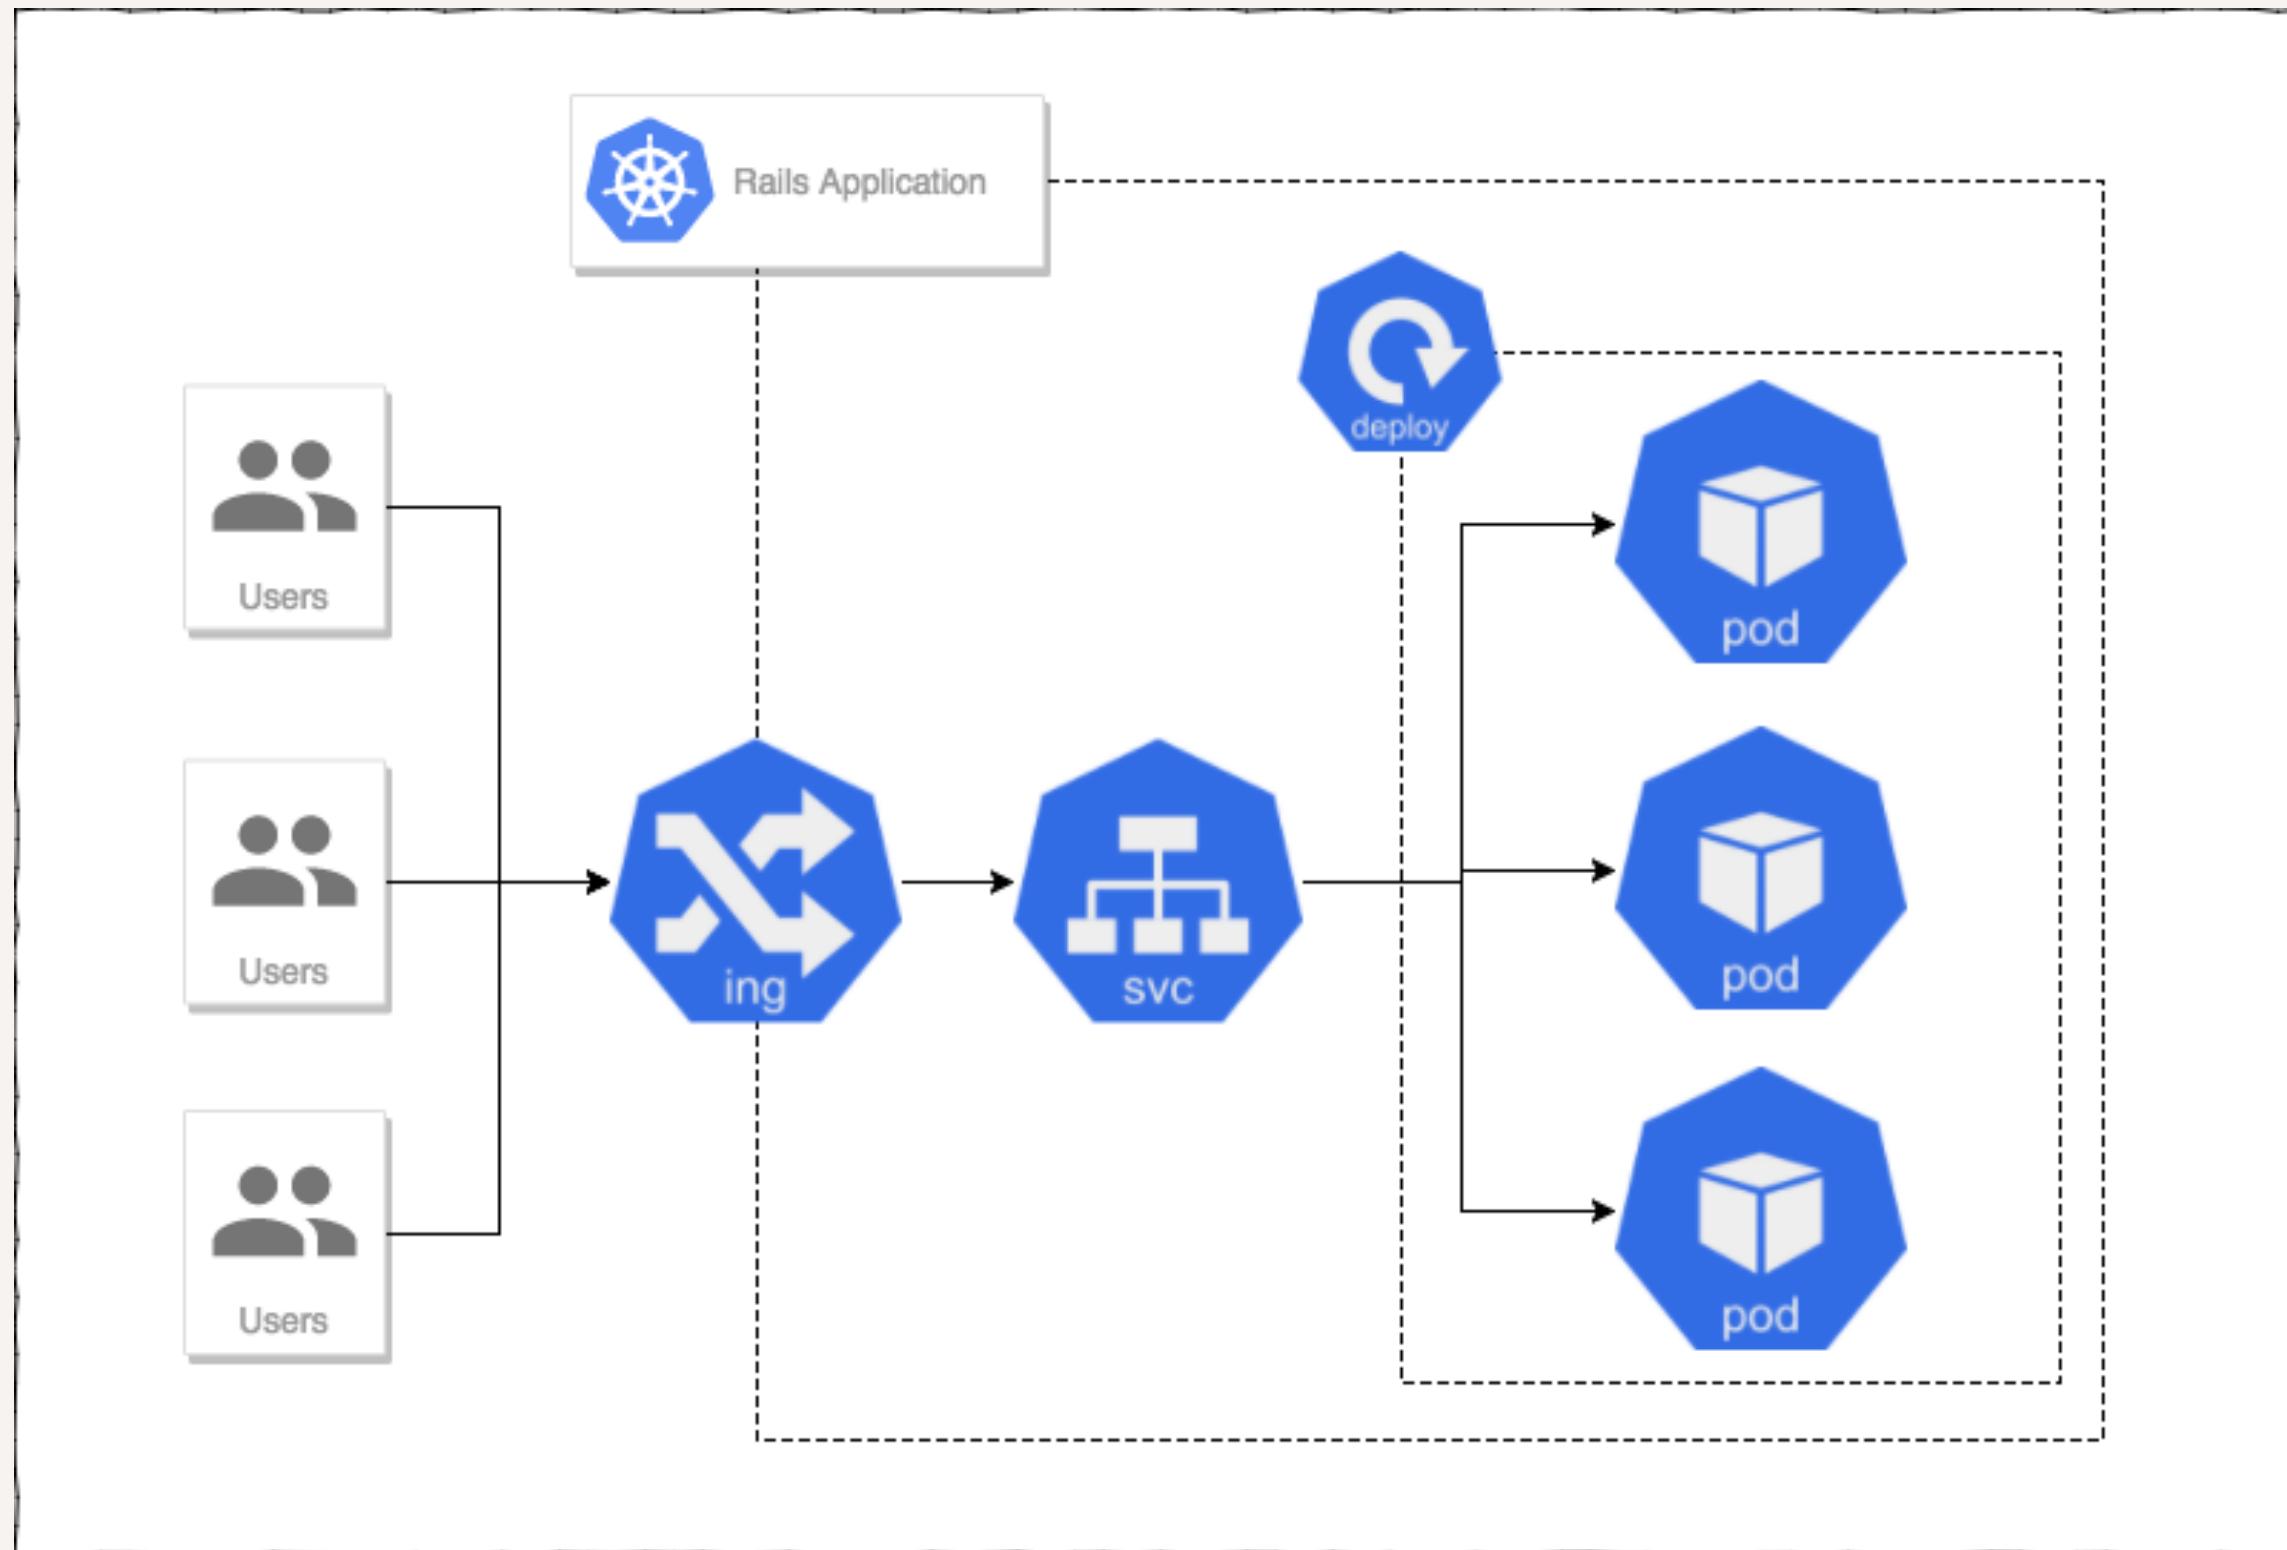
**What about talking to
external services?**



Ingress

Allows traffic to a given hostname, to route to the correct service

**In some cases, like
GKE, this even spins
up a load balancer for
you!**



```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: rails-web-ingress
spec:
  tls:
    - secretName: my-tls-cert
  rules:
    - host: my-rad-rails-app.com
      http:
        paths:
          - backend:
              serviceName: rails-web-service
              servicePort: rails-port
            path: /*
```

So lets circle back to some of our scaling problems

- Utilization and cost
- Scaling inventory with demand
- What if I need to deploy and manage more services?

**Scaling inventory
with demand**

Lets define inventory

- Servers to run our code on
- Copies of our application

**This is often called
Horizontal Scaling**

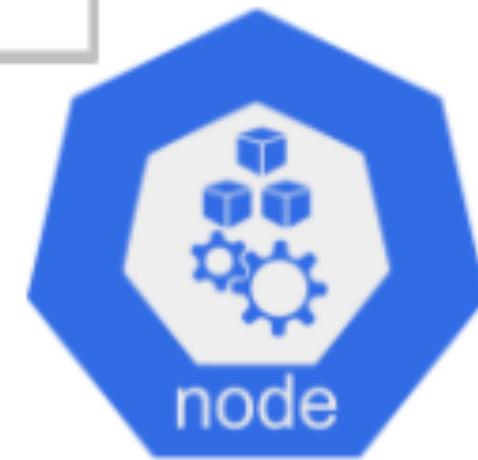
**Servers to run our
code on**

Just add more nodes!

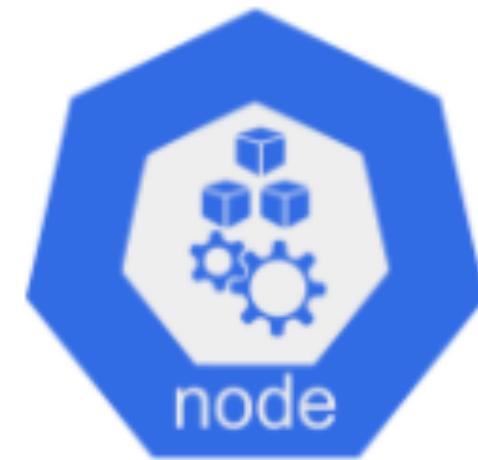
**Node can be any
mixture of machines**



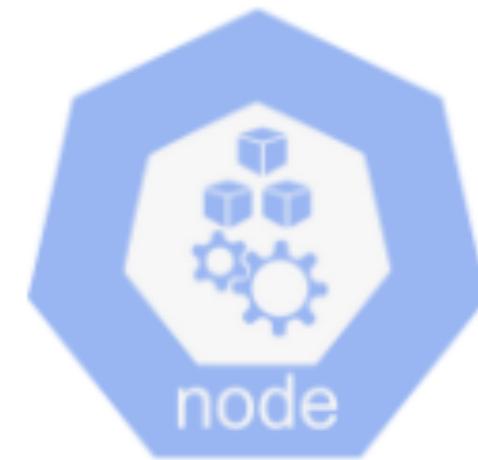
Cluster



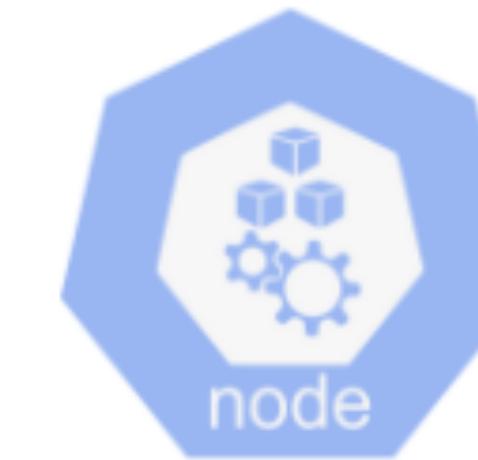
node



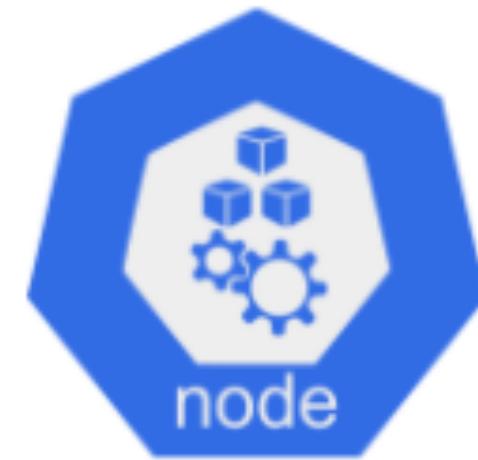
node



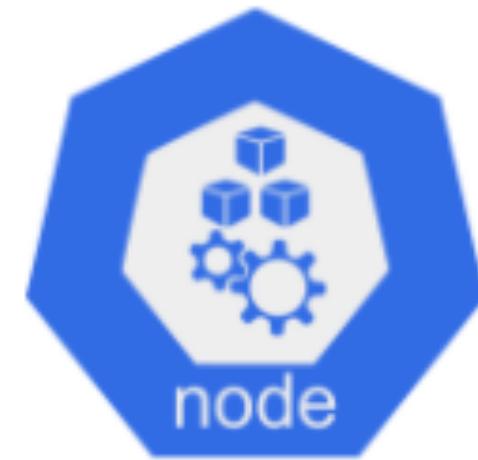
node



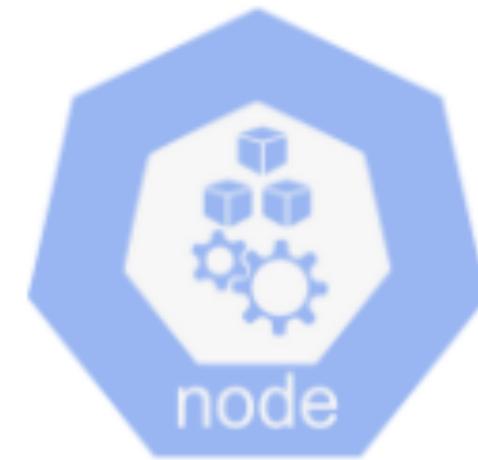
node



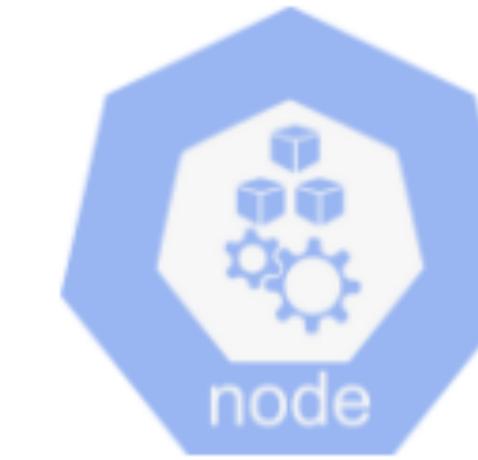
node



node



node



node

Node Autoscaler

**Scaling copies of our
application**

replicas: 3 → 4

It's that easy!

Utilization and cost

**Remember resource
requests from before?**

```
# How much resources should we allocate to this pod?  
resources:  
    requests:  
        memory: 1Gi  
        cpu: 200m  
    limits:  
        cpu: 1  
        memory: 2Gi
```

Utilization and cost

- Nodes are total *cluster* capacity
- Replicas are total *application* capacity
- These can be tweaked to find the right level of utilization

**what if I need to
deploy and manage
more services?**

A new service needs

- A docker image
- A deployment
- A service

You can then

- Deploy it separately
- Pass its service url to the Rails App
- Send requests

**Since its just a
service, you can route
traffic from inside the
cluster without
worrying about
Ingress**

**You could also try
using gRPC!**

Advanced

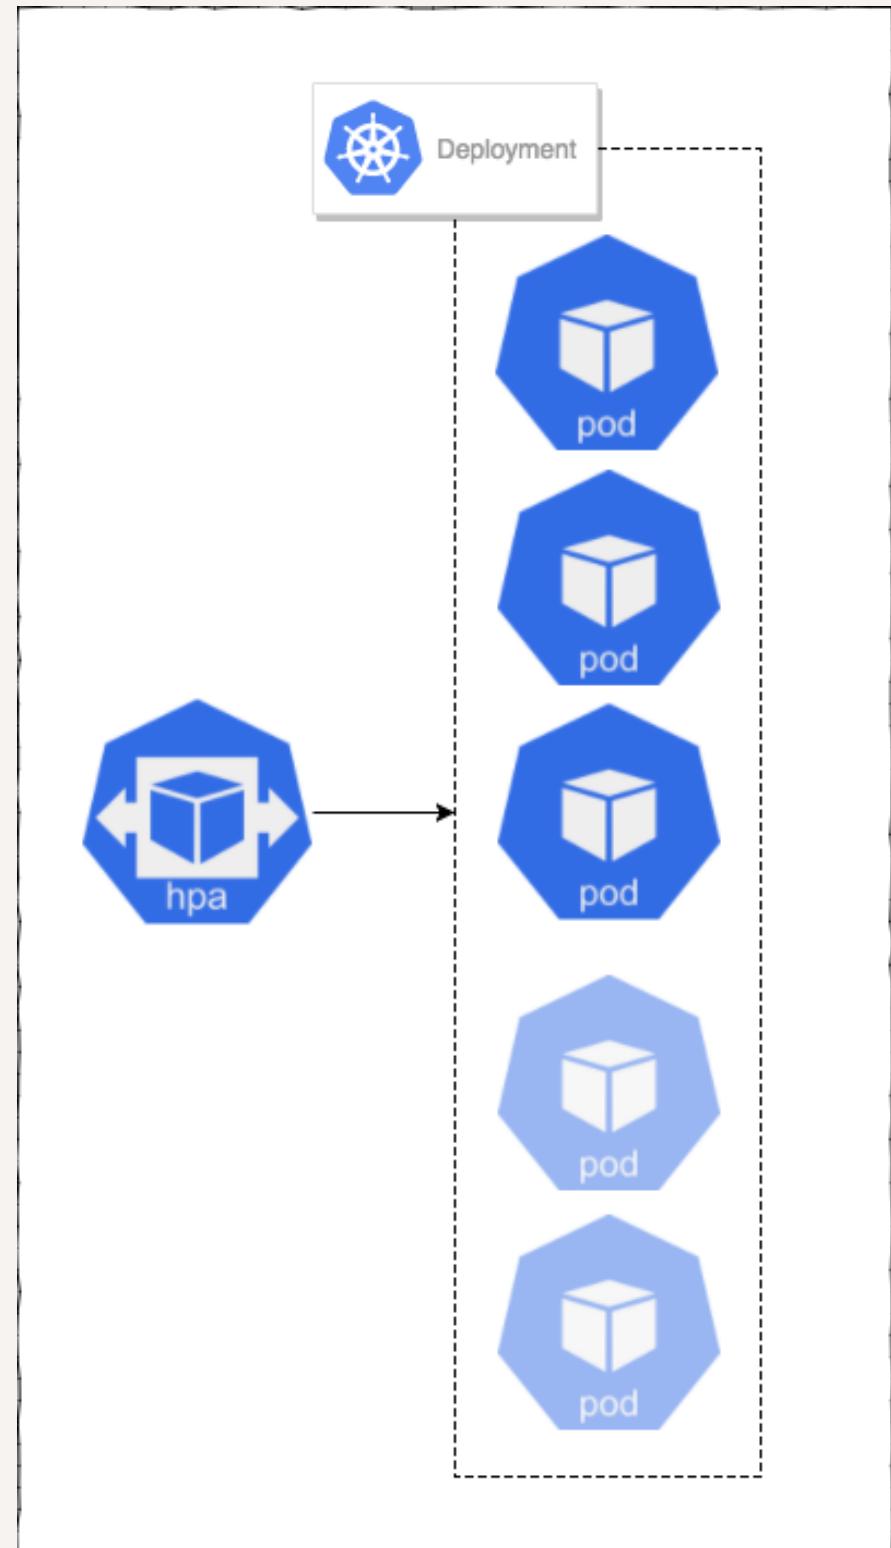
**So this is cool and all,
but show me the
good stuff!**

Autoscaling

**We already know we
can autoscale nodes,
but what about pods?**

Horizontal Pod Autoscaler

The Horizontal Pod Autoscaler automatically scales the number of pods in a replication controller, deployment or replica set based on observed CPU utilization.



```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: rails-hpa
spec:
  maxReplicas: 5
  minReplicas: 3
  scaleTargetRef:
    apiVersion: extensions/v1beta1
    kind: Deployment
    name: rails-web
  targetCPUUtilizationPercentage: 50
```

**But that's not always
the right thing to
scale on....**

Custom Metrics

**Kubernetes allows
you to extend its API**

For us,

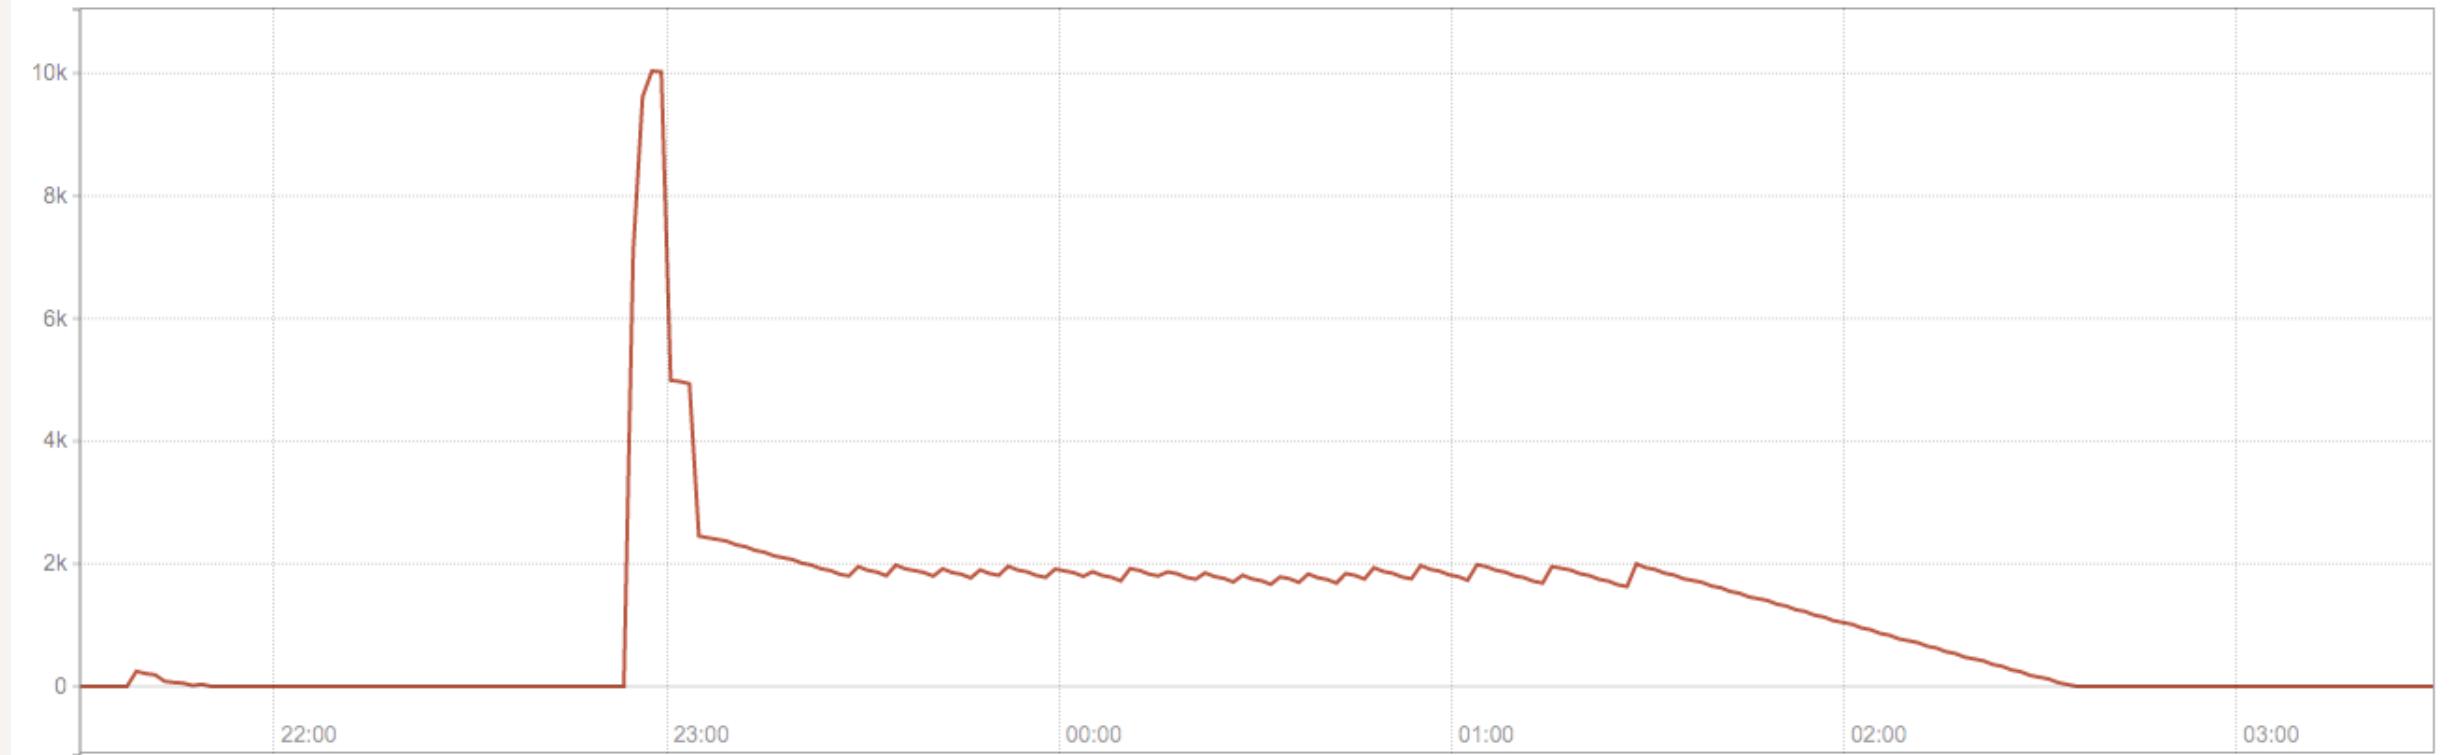
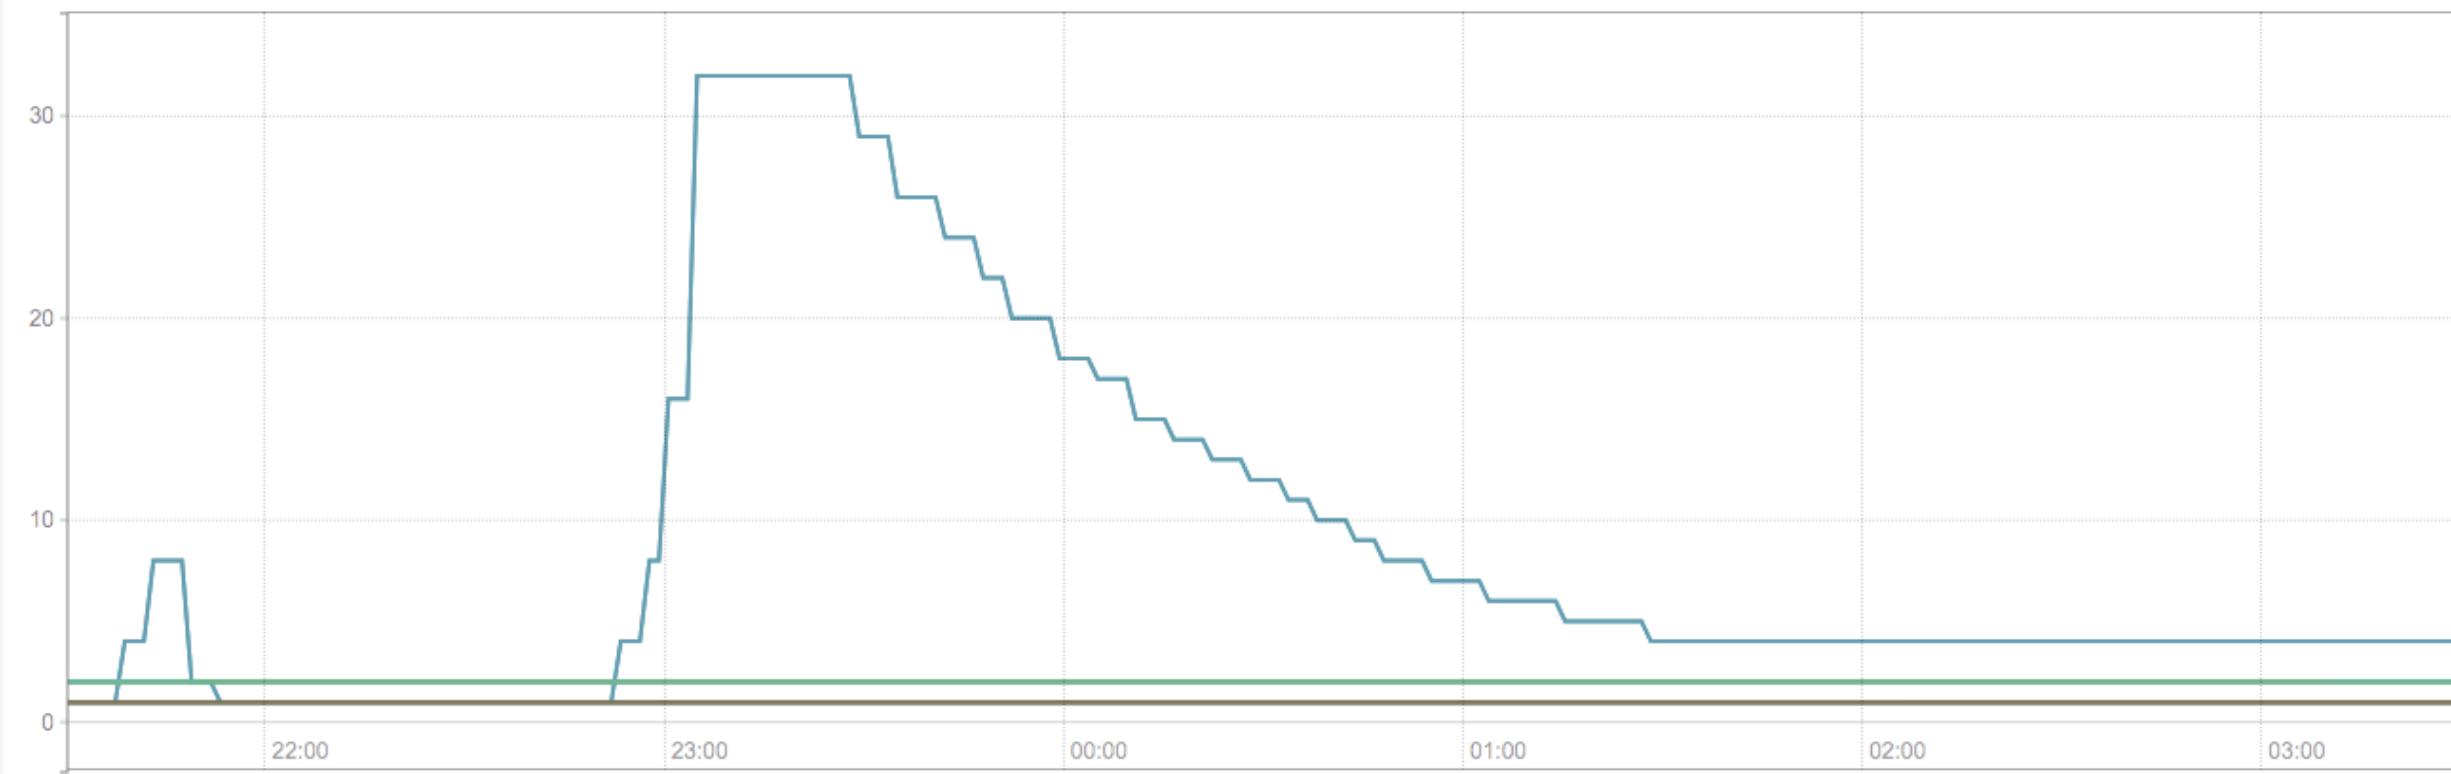
- Scaling our jobs on the `jobs_per_worker` ratio in the queue
- Scaling our web server on the `http_requests` per minute

Lets look at queues

We need to,

- Collect metrics
- Expose those metrics to Kubernetes
- Have our HPA point at those metrics





Custom Resource Definitions

**Lets think about a
worker/queue setup**

For every queue,

- I need a worker pool
- I want that worker pool to autoscale

**Kubernetes lets us
define custom
resources to create
new abstractions**

You need to...

- Define a CustomResourceDefinition
- Implement the controller that says "what do I do to fulfill this resource"

```
apiVersion: customResources/v1
kind: Queue
metadata:
  name: jira-queue
spec:
  name: jira
  min: 1
  max: 10
  metric: jobs_per_worker
```

The possibilities are endless

- A Database resource that spins up a database in cluster for you
- A Staging resource, the configures a staging app with dns behind company firewall

**This to me is the true
power of Kubernetes**

Summing up

Kubernetes facilitates..

- Thinking of infrastructure as code
- Scaling manually and automatically with ease
- Building our own abstractions

**How can I learn
more?**

- [noobernetes.io](#)
- [Tutorials - Kubernetes](#)
- [Getting started with Kubernetes as an Application Developer](#)
- [The Childrens Illustrated Guide to Kubernetes](#)

Fin

